

# beta\_vae+classifier

April 26, 2019

## 0.1 Beta -Variational AutoEncoder:

If each variable in the inferred latent representation  $z$  is only sensitive to one single generative factor and relatively invariant to other factors, we will say this representation is disentangled or factorized. One benefit that often comes with disentangled representation is good interpretability and easy generalization to a variety of tasks.

For example, a model trained on photos of human faces might capture the gentle, skin color, hair color, hair length, emotion, whether wearing a pair of glasses and many other relatively independent factors in separate dimensions. Such a disentangled representation is very beneficial to facial image generation.

-VAE (Higgins et al., 2017) is a modification of Variational Autoencoder with a special emphasis to discover disentangled latent factors. Following the same incentive in VAE, we want to maximize the probability of generating real data, while keeping the distance between the real and estimated posterior distributions small (say, under a small constant ):

1. Understanding - disentagled latent spaces
2. Using classifier - to understand its performance, VAE (detach updates) + Classifier => performance
3. VAE + classifier

```
In [0]: from glob import glob
import os
import numpy as np
import matplotlib.pyplot as plt
import shutil
from torchvision import transforms
from torchvision import models
import torchvision
import torch
from torch.autograd import Variable
import torch.nn as nn
from torch.optim import lr_scheduler
from torch import optim
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import time
from torchvision import datasets, transforms
import torch.nn.functional as F
```

```

import torch.optim as optim

import os
from torchvision.utils import save_image

## Plotting library
from matplotlib.offsetbox import OffsetImage, AnnotationBbox
from matplotlib.backends.backend_agg import FigureCanvasAgg as FigureCanvas

from scipy.stats import norm
from sklearn import manifold

%matplotlib inline

print('Torch', torch.__version__, 'CUDA', torch.version.cuda)
print('Device:', torch.device('cuda:0'))
print(torch.cuda.is_available())

is_cuda = torch.cuda.is_available()
device = torch.device ( "cuda:0" if torch.cuda.is_available () else "cpu" )

```

```

Torch 1.0.1.post2 CUDA 10.0.130
Device: cuda:0
True

```

```

In [0]: ## Show image
def imshow(img,title=None):
    """Imshow for Tensor."""
    img = img.numpy().transpose((1,2,0))
    mean = np.array([0.485, 0.456, 0.406])
    std = np.array([0.229, 0.224, 0.225])
    img = std * img + mean # normalize
    img = np.clip(img, 0, 1) # clip image
    plt.figure(figsize=(16,4))
    plt.axis('off')
    plt.imshow(img)

    if title is not None:
        plt.title(title)

def plot_grid(inputs):
    # Make a grid from batch
    out = torchvision.utils.make_grid(inputs,10,10)
    imshow(out, title="")

```

```

## Visualize some images in the dataset
def visualizeDataset(X):
    for i,image in enumerate(X):
        cv2.imshow(str(i),image)
        cv2.waitKey()
        cv2.destroyAllWindows()

def plot_loss(y, title):
    plt.figure()
    plt.plot(y)
    plt.title(title)
    plt.xlabel('epochs')
    plt.ylabel('Loss')

def plot_accuracy(y, title):
    plt.figure()
    plt.plot(y)
    plt.title(title)
    plt.xlabel('epochs')
    plt.ylabel('accuracy')

## Scatter Plot
def scatterplot(x, y, ax, imageData, zoom):
    images = []
    imageSize = 28
    for i in range(len(x)):
        x0, y0 = x[i], y[i]

        # Convert to image
        img = imageData[i]*255.
        img = (img).numpy()
        img = img.astype(np.uint8).reshape([imageSize,imageSize])
        img = cv2.cvtColor(img,cv2.COLOR_GRAY2RGB)

        # Note: OpenCV uses BGR and plt uses RGB
        image = OffsetImage(img, zoom=zoom)
        ab = AnnotationBbox(image, (x0, y0), xycoords='data', frameon=False)
        images.append(ax.add_artist(ab))

    ax.update_datalim(np.column_stack([x, y]))
    ax.autoscale()

import seaborn as sns
palette = np.array(sns.color_palette("hls", 10))

def plot_scatter(projection,labels):

```

```

plt.scatter(projection[:,0], projection[:,1],c=[palette[i] for i in labels])

In [0]: class Params:
        nb_latents = 10
        batch_size = 128
        epochs = 100
        log_interval = 100
        save_interval = 1000

        torch.manual_seed(5)

Out[0]: <torch._C.Generator at 0x7fa38ae965f0>

In [0]: """
        Traverse Latents
        """

def traverse_latents(model, datapoint, nb_latents, file):
    model.eval()
    datapoint = datapoint.to(device)
    if isinstance(model, ConvVAE):
        datapoint = datapoint.unsqueeze(0)
        mu, _ = model.encode(datapoint)
    else:
        datapoint = datapoint.unsqueeze(0)
        mu, _ = model.encoder(datapoint)

    recons = torch.zeros((7, nb_latents, 28, 28))

    for zi in range(nb_latents):
        muc = mu.squeeze().clone()
        for i, val in enumerate(np.linspace(-3, 3, 7)):
            muc[zi] = val
            recon = model.decode(muc).cpu() if isinstance(model, ConvVAE)
            else model.decoder(muc).cpu()
            recons[i, zi] = recon.view(28, 28)

    torchvision.utils.save_image(recons.view(-1, 1, 28, 28), file,
                                nrow=nb_latents, pad_value=1)

```

## 0.2 Metrics

```

In [0]: ### Metrics - Base Class For all Metrics
class Metric:
    def __init__(self):
        pass
    def __call__(self, outputs, target, loss):
        raise NotImplementedError

```

```

def reset(self):
    raise NotImplementedError

def value(self):
    raise NotImplementedError

def name(self):
    raise NotImplementedError

## Accuracy Metric
class AccumulatedAccuracyMetric(Metric):
    def __init__(self):
        self.correct = 0
        self.total = 0

    def __call__(self, outputs, target):
        # Track the accuracy
        _, argmax = torch.max(outputs, 1)
        accuracy = (target == argmax.squeeze()).float().sum()
        self.correct += accuracy
        self.total += target.size(0)
        return self.value()

    def reset(self):
        self.correct = 0
        self.total = 0

    def value(self):
        return 100 * float(self.correct) / self.total

    def name(self):
        return 'Accuracy'

## Loss
class RunningAverage():
    """A simple class that maintains the running average of a quantity
    Example:
    ...

    loss_avg = RunningAverage()
    loss_avg.update(2)
    loss_avg.update(4)
    loss_avg() = 3
    ...
    """

    def __init__( self ):
        self.steps = 0

```

```

        self.total = 0

    def update( self, val ):
        self.total += val
        self.steps += 1

    def __call__( self ):
        return self.total / float ( self.steps )

```

### 0.3 Dataloader

```

In [ ]: path = "./vae-classifier/"
        transformation = transforms.Compose([transforms.ToTensor(),
                                             transforms.Normalize((0.1307,), (0.3081,))])
        train_dataset = datasets.MNIST(os.path.join(path, "MNIST/data"),
                                         train=True, transform=transformation, download=True)
        test_dataset = datasets.MNIST(os.path.join(path, "MNIST/data"),
                                       train=False, transform=transformation, download=True)
        train_loader = torch.utils.data.DataLoader(train_dataset,
                                                    batch_size=128, shuffle=True)
        test_loader = torch.utils.data.DataLoader(test_dataset,
                                                  batch_size=128, shuffle=True)
        ## Test TSNE plot for reconstruction on 1000 test samples
        testing_tsne = torch.utils.data.DataLoader(train_dataset,
                                                    batch_size=len(train_dataset), shuffle=True)
        test_data, test_labels = next(iter(testing_tsne))[10000]

```

```

In [0]: print(f"Total number of train images: {len(train_dataset)}, " +
             f"total number of test images: {len(test_dataset)}, " +
             f"total number of train batches: {len(train_loader)}")

```

Total number of train images: 60000, total number of test images: 10000, total number of train

```

In [0]: %ls

```

```

results/  sample_data/  vae-classifier/

```

### 0.4 Model

```

In [0]: class ConvVAE(nn.Module):
        def __init__(self, nb_latents):
            super(ConvVAE, self).__init__()
            self.conv1 = nn.Sequential(
                nn.Conv2d(1, 32, kernel_size=5, stride=1, padding=2),
                nn.ReLU(),
                nn.MaxPool2d(kernel_size=2, stride=2))
            self.conv2 = nn.Sequential(

```

```

        nn.Conv2d(32, 64, kernel_size=5, stride=1, padding=2),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

self.conv3 = nn.Sequential(
    nn.Conv2d(64, 128, kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2))

self.conv4 = nn.Sequential(
    nn.Conv2d(128, 256, kernel_size=2, stride=1),
    nn.ReLU())

self.fc1 = nn.Linear(1024, 256)

self.fc_mean = nn.Linear(256, nb_latents)
self.fc_std = nn.Linear(256, nb_latents)

self.fc2 = nn.Linear(nb_latents, 256)
self.fc3 = nn.Linear(256, 1024)

self.fc4 = nn.Linear(1024, 7*7*64)

self.deconv1 = nn.ConvTranspose2d(64, 32,
                                   kernel_size=2, stride=2)
self.deconv2 = nn.ConvTranspose2d(32, 1,
                                   kernel_size=2, stride=2)

self.relu = nn.ReLU()
self.sigmoid = nn.Sigmoid()

def encode(self, x):
    x = (self.conv1(x))
    x = (self.conv2(x))
    x = (self.conv3(x))
    x = (self.conv4(x))
    x = x.reshape(x.size(0), -1)
    x = self.relu(self.fc1(x))
    return self.fc_mean(x), self.fc_std(x)

def reparameterize(self, mu, logvar):
    if self.training:
        std = logvar.mul(0.5).exp_()
        eps = Variable(std.data.new(std.size()).normal_())
        return eps.mul(std).add_(mu)

```

```

        else:
            return mu

    def decode(self, z):
        x = self.relu(self.fc2(z))
        x = self.relu(self.fc3(x))
        x = self.relu(self.fc4(x))
        x = self.relu(self.deconv1(x.view(-1, 64, 7, 7)))
        x = self.deconv2(x)
        return self.sigmoid(x)

    def forward(self, x):
        mu, logvar = self.encode(x)
        z = self.reparameterize(mu, logvar)
        return self.decode(z), mu, logvar

In [0]: model = ConvVAE(Params.nb_latents)
        if is_cuda:
            model = model.to(device)

        print(model)

ConvVAE(
  (conv1): Sequential(
    (0): Conv2d(1, 32, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (conv2): Sequential(
    (0): Conv2d(32, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (conv3): Sequential(
    (0): Conv2d(64, 128, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (1): ReLU()
    (2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (conv4): Sequential(
    (0): Conv2d(128, 256, kernel_size=(2, 2), stride=(1, 1))
    (1): ReLU()
  )
  (fc1): Linear(in_features=1024, out_features=256, bias=True)
  (fc_mean): Linear(in_features=256, out_features=10, bias=True)
  (fc_std): Linear(in_features=256, out_features=10, bias=True)
  (fc2): Linear(in_features=10, out_features=256, bias=True)
  (fc3): Linear(in_features=256, out_features=1024, bias=True)
  (fc4): Linear(in_features=1024, out_features=3136, bias=True)

```



```

(deconv1): ConvTranspose2d(64, 32, kernel_size=(2, 2), stride=(2, 2))
(deconv2): ConvTranspose2d(32, 1, kernel_size=(2, 2), stride=(2, 2))
(relu): ReLU()
(sigmoid): Sigmoid()
)

```

## 0.5 Loss Function

```

In [0]: ### Loss Function
        ### Loss Function
        def loss_function(recon_x, x, mu, logvar, beta=6):
            """
            Reconstruction loss + KL divergence loss over all elements of the batch
            """

            bce = F.binary_cross_entropy(recon_x,
                                         x.view(-1, 28*28), size_average=False)

            kld = -0.5* (1+ logvar -mu.pow(2) - logvar.exp())
            return kld.mean(dim = 0), bce + beta*kld.sum()

In [0]: import os
        os.makedirs("beta-results")
        %ls

        save_dir = "beta-results"

beta-results/  results/  sample_data/  vae-classifier/

```

## 0.6 Latent Space Visualization

```

In [0]: criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(model.parameters(), lr=1e-3)

        """
        Traverse Latents
        """

        def traverse_latents(model, datapoint, nb_latents,
                              epoch, batch_idx, dirpath=save_dir):

            model.eval()
            datapoint = datapoint.to(device)
            if isinstance(model, ConvVAE):
                datapoint = datapoint.unsqueeze(0)
                mu, _ = model.encode(datapoint)
            else:
                mu, _ = model.encode(datapoint.view(-1))

```

```

recons = torch.zeros((7, nb_latents, 28, 28))
for zi in range(nb_latents):
    muc = mu.squeeze().clone()
    for i, val in enumerate(np.linspace(-3, 3, 7)):
        muc[zi] = val

    recon = model.decode(muc).cpu()
    recons[i, zi] = recon.view(28, 28)

filename = os.path.join(dirpath, 'traversal_' +
                        str(epoch) + '_' + str(batch_idx) + '.png')
save_image(recons.view(-1, 1, 28, 28),
           filename, nrow=nb_latents, pad_value=1)

### Plot TSNE for latent space
# Show dataset images with T-sne projection of latent space encoding
import seaborn as sns
palette = np.array(sns.color_palette("hls", 10))
def visualize_tsne_of_input(X, labels, model, path):
    # Compute latent space representation
    print("Computing latent space projection...")

    X_encoded, _ = model.encode(X)

    # Compute t-SNE embedding of latent space
    tsne = manifold.TSNE(n_components=2, init='pca', random_state=0)
    X_tsne = tsne.fit_transform(X_encoded.data.detach().cpu())

    # Plot images according to t-sne embedding
    fig, ax = plt.subplots()
    im = plt.scatter(X_tsne[:,0], X_tsne[:,1],
                    c=[palette[i] for i in labels])
    fig.colorbar(im, ticks=range(10));
    fig.savefig(path, dpi=fig.dpi)

```

```

In [0]: #!pip install tqdm
        !pip install opencv-python

```

Requirement already satisfied: opencv-python in /home/chivukula\_manju/yes/lib/python3.6/site-packages  
Requirement already satisfied: numpy>=1.11.3 in /home/chivukula\_manju/yes/lib/python3.6/site-packages  
You are using pip version 9.0.1, however version 19.0.3 is available.You should consider upgrading

## 0.7 Compute Laplacian Variance for the blur - cv2

```

In [0]: ## Import CV
        import cv2

```

```

import numpy

def laplacian_variance(images):
    return [cv2.Laplacian(image.numpy(), cv2.CV_32F).var() for image in images]

def laplacian_variance_numpy(images):
    return [cv2.Laplacian(image, cv2.CV_32F).var() for image in images]

In [0]: log_interval = 400
        save_interval = 400
        testpoint = torch.Tensor(train_loader.dataset[0][0]).to(device)

def train_epoch(epoch):
    model.train()
    train_losses = RunningAverage()
    for batch_idx, (data, labels) in enumerate(train_loader):
        data = data.to(device)
        labels = labels.to(device)
        batch_size = data.size(0)
        recon_batch, mu, logvar = model(data)

        kld, loss = loss_function(recon_batch.squeeze()
                                  .view(-1, 28*28), data, mu, logvar)
        ## parameter update
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        loss /= len(data)
        train_losses.update(loss.item())
    return train_losses()

In [0]: ## Test Epoch
        """
        Test
        """

def test_epoch(epoch):
    model.eval()
    test_loss = RunningAverage()

    ## Test reconstruction test
    testpoint = torch.Tensor(test_loader.dataset[0][0]).to(device)

    with torch.no_grad():
        for i, (data, labels) in enumerate(test_loader):

```

```

        data = data.to(device)

        batch_size = data.size(0)
        recon_batch, mu, logvar = model(data)

        ## Loss of VAE
        kld, loss = loss_function(recon_batch
                                   .squeeze().view(-1,28*28), data, mu, logvar)
        test_loss.update(loss.item() / len(data))

    return test_loss()

In [0]: train_losses = []
        test_losses = []
        for epoch in range(Params.epochs):
            train_loss = train_epoch(epoch)
            train_losses.append(train_loss)
            message = 'Epoch: {}/{}. Train set: ' +
                      'Average loss: {:.4f}'.format(epoch + 1,
                                                    Params.epochs, train_loss)
            print(message)
            test_loss = test_epoch(epoch)
            message = 'Epoch: {}/{}. Test set: ' +
                      'Average loss: {:.4f}'.format(epoch + 1,
                                                    Params.epochs, test_loss)
            test_losses.append(test_loss)
            print(message)

```

```

/home/chivukula_manju/yes/lib/python3.6/site-packages/torch/nn/_reduction.py:49: UserWarning: s
warnings.warn(warning.format(ret))

```

```

Epoch: 1/100. Train set: Average loss: -10349.4857
Epoch: 1/100. Test set: Average loss: -10544.8887
Epoch: 2/100. Train set: Average loss: -10452.9163
Epoch: 2/100. Test set: Average loss: -10587.9207
Epoch: 3/100. Train set: Average loss: -10512.6402
Epoch: 3/100. Test set: Average loss: -10684.3192
Epoch: 4/100. Train set: Average loss: -10562.4123
Epoch: 4/100. Test set: Average loss: -10696.9424
Epoch: 5/100. Train set: Average loss: -10611.1751
Epoch: 5/100. Test set: Average loss: -10748.8865
Epoch: 6/100. Train set: Average loss: -10661.2261
Epoch: 6/100. Test set: Average loss: -10813.1165
Epoch: 7/100. Train set: Average loss: -10708.5552
Epoch: 7/100. Test set: Average loss: -10825.0708

```

Epoch: 8/100. Train set: Average loss: -10744.1072  
Epoch: 8/100. Test set: Average loss: -10855.0397  
Epoch: 9/100. Train set: Average loss: -10763.4951  
Epoch: 9/100. Test set: Average loss: -10909.9545  
Epoch: 10/100. Train set: Average loss: -10783.4877  
Epoch: 10/100. Test set: Average loss: -10902.9798  
Epoch: 11/100. Train set: Average loss: -10799.0980  
Epoch: 11/100. Test set: Average loss: -10885.9479  
Epoch: 12/100. Train set: Average loss: -10813.8132  
Epoch: 12/100. Test set: Average loss: -10896.8893  
Epoch: 13/100. Train set: Average loss: -10830.2422  
Epoch: 13/100. Test set: Average loss: -10931.7071  
Epoch: 14/100. Train set: Average loss: -10835.9709  
Epoch: 14/100. Test set: Average loss: -10937.8951  
Epoch: 15/100. Train set: Average loss: -10852.6770  
Epoch: 15/100. Test set: Average loss: -10942.0391  
Epoch: 16/100. Train set: Average loss: -10860.5345  
Epoch: 16/100. Test set: Average loss: -10955.7746  
Epoch: 17/100. Train set: Average loss: -10869.8607  
Epoch: 17/100. Test set: Average loss: -10967.9799  
Epoch: 18/100. Train set: Average loss: -10874.4623  
Epoch: 18/100. Test set: Average loss: -10985.3355  
Epoch: 19/100. Train set: Average loss: -10890.4998  
Epoch: 19/100. Test set: Average loss: -10977.6311  
Epoch: 20/100. Train set: Average loss: -10899.1664  
Epoch: 20/100. Test set: Average loss: -10981.4035  
Epoch: 21/100. Train set: Average loss: -10903.4473  
Epoch: 21/100. Test set: Average loss: -10978.4739  
Epoch: 22/100. Train set: Average loss: -10912.3935  
Epoch: 22/100. Test set: Average loss: -10983.2029  
Epoch: 23/100. Train set: Average loss: -10915.2949  
Epoch: 23/100. Test set: Average loss: -10967.9590  
Epoch: 24/100. Train set: Average loss: -10918.3851  
Epoch: 24/100. Test set: Average loss: -10977.4329  
Epoch: 25/100. Train set: Average loss: -10925.6430  
Epoch: 25/100. Test set: Average loss: -10960.0561  
Epoch: 26/100. Train set: Average loss: -10930.3654  
Epoch: 26/100. Test set: Average loss: -10987.0091  
Epoch: 27/100. Train set: Average loss: -10930.2887  
Epoch: 27/100. Test set: Average loss: -11004.4141  
Epoch: 28/100. Train set: Average loss: -10935.9621  
Epoch: 28/100. Test set: Average loss: -11006.9101  
Epoch: 29/100. Train set: Average loss: -10939.9603  
Epoch: 29/100. Test set: Average loss: -10992.1838  
Epoch: 30/100. Train set: Average loss: -10946.9011  
Epoch: 30/100. Test set: Average loss: -10975.8259  
Epoch: 31/100. Train set: Average loss: -10949.9610  
Epoch: 31/100. Test set: Average loss: -11004.9979

Epoch: 32/100. Train set: Average loss: -10952.7151  
Epoch: 32/100. Test set: Average loss: -11015.1051  
Epoch: 33/100. Train set: Average loss: -10959.3073  
Epoch: 33/100. Test set: Average loss: -10982.9964  
Epoch: 34/100. Train set: Average loss: -10961.4371  
Epoch: 34/100. Test set: Average loss: -11008.3863  
Epoch: 35/100. Train set: Average loss: -10963.2451  
Epoch: 35/100. Test set: Average loss: -11008.2428  
Epoch: 36/100. Train set: Average loss: -10966.7791  
Epoch: 36/100. Test set: Average loss: -11011.4464  
Epoch: 37/100. Train set: Average loss: -10969.2139  
Epoch: 37/100. Test set: Average loss: -11000.2670  
Epoch: 38/100. Train set: Average loss: -10973.4729  
Epoch: 38/100. Test set: Average loss: -11013.9935  
Epoch: 39/100. Train set: Average loss: -10974.0112  
Epoch: 39/100. Test set: Average loss: -11007.8988  
Epoch: 40/100. Train set: Average loss: -10975.1199  
Epoch: 40/100. Test set: Average loss: -11001.5884  
Epoch: 41/100. Train set: Average loss: -10978.4597  
Epoch: 41/100. Test set: Average loss: -11031.9370  
Epoch: 42/100. Train set: Average loss: -10984.2374  
Epoch: 42/100. Test set: Average loss: -10990.9801  
Epoch: 43/100. Train set: Average loss: -10986.2278  
Epoch: 43/100. Test set: Average loss: -11009.4645  
Epoch: 44/100. Train set: Average loss: -10986.6053  
Epoch: 44/100. Test set: Average loss: -11016.3920  
Epoch: 45/100. Train set: Average loss: -10988.7983  
Epoch: 45/100. Test set: Average loss: -11015.9668  
Epoch: 46/100. Train set: Average loss: -10992.2750  
Epoch: 46/100. Test set: Average loss: -10995.0556  
Epoch: 47/100. Train set: Average loss: -10992.0277  
Epoch: 47/100. Test set: Average loss: -10983.6734  
Epoch: 48/100. Train set: Average loss: -10994.8473  
Epoch: 48/100. Test set: Average loss: -11002.2136  
Epoch: 49/100. Train set: Average loss: -10995.0274  
Epoch: 49/100. Test set: Average loss: -11023.6951  
Epoch: 50/100. Train set: Average loss: -10997.7335  
Epoch: 50/100. Test set: Average loss: -11019.6370  
Epoch: 51/100. Train set: Average loss: -11002.3454  
Epoch: 51/100. Test set: Average loss: -11006.3942  
Epoch: 52/100. Train set: Average loss: -11001.4557  
Epoch: 52/100. Test set: Average loss: -11022.6608  
Epoch: 53/100. Train set: Average loss: -11003.6703  
Epoch: 53/100. Test set: Average loss: -11003.6906  
Epoch: 54/100. Train set: Average loss: -11005.1104  
Epoch: 54/100. Test set: Average loss: -10993.1252  
Epoch: 55/100. Train set: Average loss: -11006.1172  
Epoch: 55/100. Test set: Average loss: -10999.5482

Epoch: 56/100. Train set: Average loss: -11008.0077  
Epoch: 56/100. Test set: Average loss: -11021.4441  
Epoch: 57/100. Train set: Average loss: -11008.7850  
Epoch: 57/100. Test set: Average loss: -11020.0875  
Epoch: 58/100. Train set: Average loss: -11011.2379  
Epoch: 58/100. Test set: Average loss: -11015.5271  
Epoch: 59/100. Train set: Average loss: -11013.7164  
Epoch: 59/100. Test set: Average loss: -11026.9246  
Epoch: 60/100. Train set: Average loss: -11012.7646  
Epoch: 60/100. Test set: Average loss: -11012.6244  
Epoch: 61/100. Train set: Average loss: -11013.7034  
Epoch: 61/100. Test set: Average loss: -11010.1708  
Epoch: 62/100. Train set: Average loss: -11016.8086  
Epoch: 62/100. Test set: Average loss: -11031.8402  
Epoch: 63/100. Train set: Average loss: -11017.4972  
Epoch: 63/100. Test set: Average loss: -11021.7985  
Epoch: 64/100. Train set: Average loss: -11016.6438  
Epoch: 64/100. Test set: Average loss: -10994.2679  
Epoch: 65/100. Train set: Average loss: -11019.7879  
Epoch: 65/100. Test set: Average loss: -11008.0334  
Epoch: 66/100. Train set: Average loss: -11018.8163  
Epoch: 66/100. Test set: Average loss: -11001.5700  
Epoch: 67/100. Train set: Average loss: -11021.1522  
Epoch: 67/100. Test set: Average loss: -11028.0907  
Epoch: 68/100. Train set: Average loss: -11022.9235  
Epoch: 68/100. Test set: Average loss: -11006.5722  
Epoch: 69/100. Train set: Average loss: -11023.1485  
Epoch: 69/100. Test set: Average loss: -11008.3074  
Epoch: 70/100. Train set: Average loss: -11024.1373  
Epoch: 70/100. Test set: Average loss: -11010.4841  
Epoch: 71/100. Train set: Average loss: -11026.2412  
Epoch: 71/100. Test set: Average loss: -11011.2571  
Epoch: 72/100. Train set: Average loss: -11027.5816  
Epoch: 72/100. Test set: Average loss: -11031.9309  
Epoch: 73/100. Train set: Average loss: -11029.1204  
Epoch: 73/100. Test set: Average loss: -10991.2853  
Epoch: 74/100. Train set: Average loss: -11026.8116  
Epoch: 74/100. Test set: Average loss: -11024.3664  
Epoch: 75/100. Train set: Average loss: -11029.0669  
Epoch: 75/100. Test set: Average loss: -10992.8561  
Epoch: 76/100. Train set: Average loss: -11031.9145  
Epoch: 76/100. Test set: Average loss: -11003.8241  
Epoch: 77/100. Train set: Average loss: -11032.3702  
Epoch: 77/100. Test set: Average loss: -11022.3059  
Epoch: 78/100. Train set: Average loss: -11032.2286  
Epoch: 78/100. Test set: Average loss: -11029.7318  
Epoch: 79/100. Train set: Average loss: -11034.1605  
Epoch: 79/100. Test set: Average loss: -11016.2777

Epoch: 80/100. Train set: Average loss: -11036.1964  
Epoch: 80/100. Test set: Average loss: -11026.5891  
Epoch: 81/100. Train set: Average loss: -11036.9252  
Epoch: 81/100. Test set: Average loss: -11008.2321  
Epoch: 82/100. Train set: Average loss: -11037.5574  
Epoch: 82/100. Test set: Average loss: -10998.7115  
Epoch: 83/100. Train set: Average loss: -11036.8492  
Epoch: 83/100. Test set: Average loss: -11009.0513  
Epoch: 84/100. Train set: Average loss: -11039.0210  
Epoch: 84/100. Test set: Average loss: -11020.4130  
Epoch: 85/100. Train set: Average loss: -11038.9508  
Epoch: 85/100. Test set: Average loss: -11016.6073  
Epoch: 86/100. Train set: Average loss: -11039.4227  
Epoch: 86/100. Test set: Average loss: -11011.6423  
Epoch: 87/100. Train set: Average loss: -11040.1777  
Epoch: 87/100. Test set: Average loss: -11015.3848  
Epoch: 88/100. Train set: Average loss: -11038.0958  
Epoch: 88/100. Test set: Average loss: -11004.8606  
Epoch: 89/100. Train set: Average loss: -11040.6147  
Epoch: 89/100. Test set: Average loss: -11020.3086  
Epoch: 90/100. Train set: Average loss: -11041.2143  
Epoch: 90/100. Test set: Average loss: -11002.2994  
Epoch: 91/100. Train set: Average loss: -11043.8785  
Epoch: 91/100. Test set: Average loss: -11019.4219  
Epoch: 92/100. Train set: Average loss: -11045.6124  
Epoch: 92/100. Test set: Average loss: -10999.1384  
Epoch: 93/100. Train set: Average loss: -11043.5832  
Epoch: 93/100. Test set: Average loss: -11004.1811  
Epoch: 94/100. Train set: Average loss: -11043.8821  
Epoch: 94/100. Test set: Average loss: -10998.8833  
Epoch: 95/100. Train set: Average loss: -11046.3663  
Epoch: 95/100. Test set: Average loss: -11007.7211  
Epoch: 96/100. Train set: Average loss: -11046.6947  
Epoch: 96/100. Test set: Average loss: -11008.7210  
Epoch: 97/100. Train set: Average loss: -11047.7590  
Epoch: 97/100. Test set: Average loss: -10996.8445  
Epoch: 98/100. Train set: Average loss: -11046.2223  
Epoch: 98/100. Test set: Average loss: -11001.2824  
Epoch: 99/100. Train set: Average loss: -11047.3682  
Epoch: 99/100. Test set: Average loss: -11009.4789  
Epoch: 100/100. Train set: Average loss: -11047.8037  
Epoch: 100/100. Test set: Average loss: -11005.8705

-----  
TypeError

Traceback (most recent call last)



```

<ipython-input-49-1f4fefb2aaed> in <module>()
    12
    13
---> 14 traverse_latents(model, Params.nb_latents, epoch,save_dir)
    15 plt.plot(range(Params.epochs),train_losses)
    16 plt.plot(range(Params.epochs),test_losses)

```

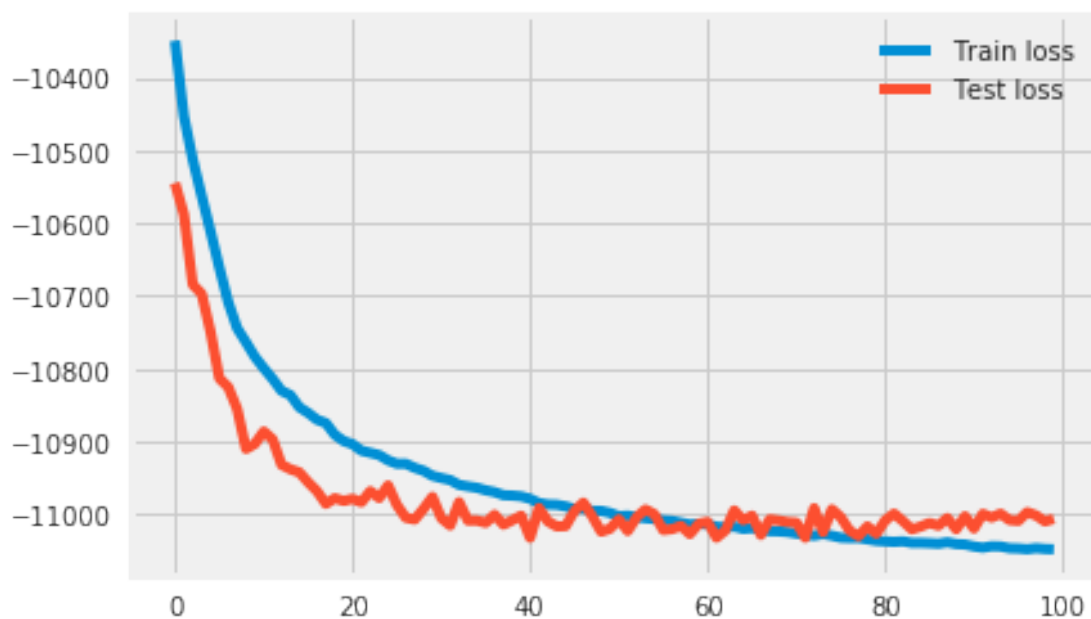
TypeError: traverse\_latents() missing 1 required positional argument: 'batch\_idx'

```

In [0]: # Plot Losses
plt.plot(range(Params.epochs),train_losses)
plt.plot(range(Params.epochs),test_losses)
plt.legend(["Train loss","Test loss"])
plt.show()

```

/home/chivukula\_manju/yes/lib/python3.6/site-packages/matplotlib/font\_manager.py:1328: UserWarning  
(prop.get\_family(), self.defaultFamily[fonttext]))



```

In [0]: ## save_model
## Save model
def save_model(model,path):

    torch.save(model.state_dict(),path)

```

```
def load_model(model,path):
    model.load_state_dict(torch.load(path))

path = save_dir+ "/model_state_" + str(epoch) + ".pth"
save_model(model,path)
```

## 0.8 Reconstructions

```
In [0]: epoch = 99
path = save_dir+ "/model_state_" + str(epoch) + ".pth"

model.load_state_dict(torch.load(path))

In [0]: def reconstruction(data,epoch, is_train = True):
    n = min(data.size(0), 8)
    recon_batch, mu, logvar = model(data)
    comparison = torch.cat([data[:n],recon_batch[:n]])
    if is_train:
        name = save_dir + "/train_reconstruction_" + str(epoch) + '.png'
    else:
        name = save_dir + "/test_reconstruction_"
        + str(epoch) + '.png'
    save_image(comparison.cpu(),name, nrow=n)

train_batch, train_label = next(iter(train_loader))
test_batch, test_label = next(iter(test_loader))
reconstruction(train_batch.to(device), epoch, True)
reconstruction(test_batch.to(device), epoch, False)
```


```
In [0]: %pwd
```

```
Out[0]: '/home/chivukula_manju/vae'
```

### 0.8.1 Train Reconstruction

```
In [0]: from IPython.display import Image
Image(filename=save_dir+"/train_reconstruction_99.png")
```


```
Out[0]:
```



## 0.8.2 Test Reconstruction

```
In [0]: Image(filename=save_dir+'/test_reconstruction_99.png')
```

Out[0]:



## 0.9 Generated Samples

```
In [0]: # Generate samples of reconstruction
with torch.no_grad():
    sample = torch.randn(64, Params.nb_latents).to(device)
    sample = model.decode(sample).cpu()
    save_image(sample.view(64, 1, 28, 28),
                save_dir + '/sample_' + str(epoch) + '.png')
```

```
In [0]: from IPython.display import Image
Image(filename=save_dir+'/sample_99.png')
```

Out[0]:




## 0.10 Latent Traversals

```
In [0]: testpoint = torch.Tensor(train_loader.dataset[0][0]).to(device)
        traverse_latents(model, testpoint,
                        Params.nb_latents, epoch,1,save_dir)
```

```
In [0]: Image(filename=save_dir+'/traversal_99_1.png')
```

Out[0]:



In [0]:

## 0.11 TSNE Plot

```
In [0]: path = save_dir+'/latent_space.png'
        def visualize_tsne(X, labels, model, path):
            # Compute latent space representation
            print("Computing latent space projection...")

            X_encoded, _ = model.encode(X)

            # Compute t-SNE embedding of latent space
            tsne = manifold.TSNE(n_components=2, init='pca',
                                random_state=0)
            X_tsne = tsne.fit_transform(X_encoded.data.detach().cpu())

            # Plot images according to t-sne embedding
            fig, ax = plt.subplots()
            im = plt.scatter(X_tsne[:,0], X_tsne[:,1],c=[palette[i]
```

```

        for i in labels])

    fig.savefig(path, dpi=fig.dpi)

    visualize_tsne(test_data[:5000].to(device),
                  test_labels[:5000].to(device), model, path)


```

Computing latent space projection...

```

/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning:
  (prop.get_family(), self.defaultFamily[fonttext]))

```



## 0.12 Load model

```


In [0]: load_model(model, "./vae-classifier/results/model_state_99.pth")

In [0]: x_test, y_test = test_data[:2000], test_labels[:2000]
        laplacian_variances = [laplacian_variance(
            x_test[y_test == i]) for i in range(10)]

In [0]: plt.boxplot(laplacian_variances, labels=range(10));
        plt.xlabel('Digit class')
        plt.ylabel('Laplacian variance')
        plt.title('Laplacian variance statistics per digit class');

```

```
/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning:
(prop.get_family(), self.defaultFamily[fonttext]))
```




```
In [0]: with torch.no_grad():
        reconstructions = np.empty(shape=(227,1,28,28))
        indx = 0
        for i, (x,y) in enumerate(zip(x_test,y_test)):
            if y == 1:
                recon_batch, mu, logvar = model(x.unsqueeze(1).to(device))
                reconstructions[indx]=
                (recon_batch.squeeze(0).detach().cpu())
                indx+=1
```

```
In [0]: plt.imshow(np.array(reconstructions)[0].squeeze())
```

```
Out[0]: <matplotlib.image.AxesImage at 0x7fdebee7b470>
```

```
/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning:
(prop.get_family(), self.defaultFamily[fonttext]))
```



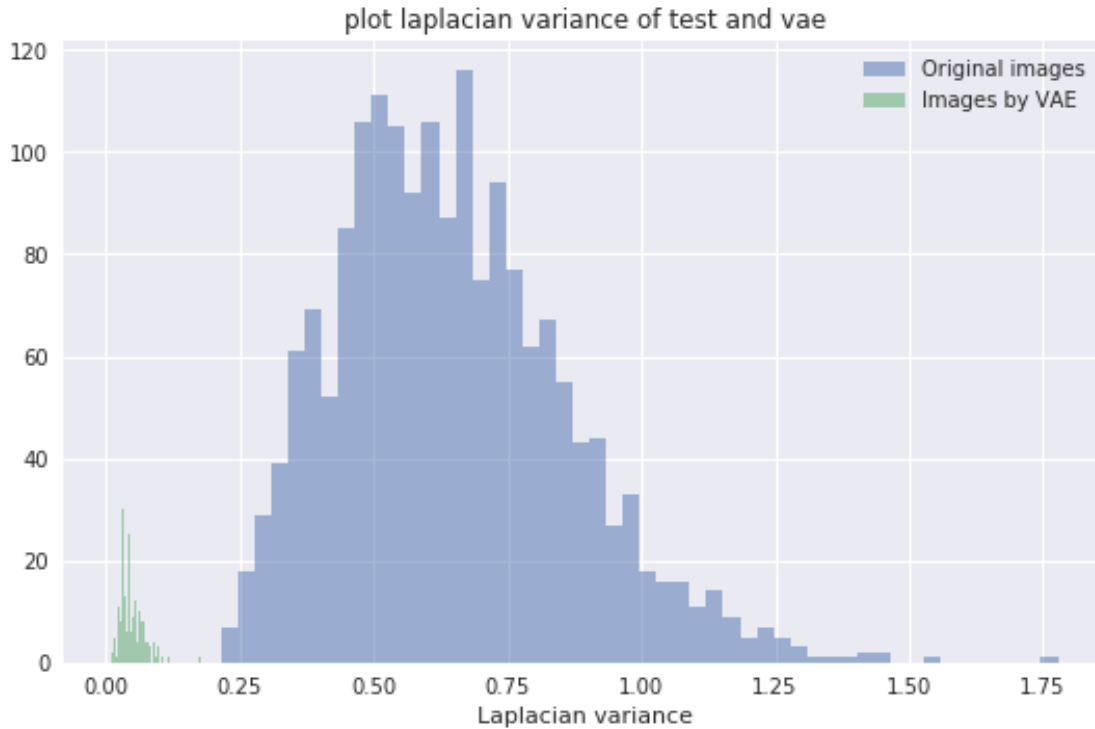
```
In [0]: laplacian_variances_recons = [laplacian_variance_numpy
                                         (np.array(reconstructions[y_test == i], dtype=np.uint8))
                                         for i in range(10)]
```

```
In [0]: not_ones = y_test != 1
```

```
lvs_1 = laplacian_variance(x_test[not_ones])
lvs_2 = laplacian_variance_numpy(np.array
                                   (reconstructions, dtype=np.uint8))
def plot_laplacian_variances(lvs_1, lvs_2, title):
    plt.hist(lvs_1, bins=50, alpha=0.5, label='Original images');
    plt.hist(lvs_2, bins=50, alpha= 0.5, label='Images by VAE');
    plt.xlabel('Laplacian variance')
    plt.title(title)
    plt.legend();

    plot_laplacian_variances(lvs_1, lvs_2, "plot laplacian variance of test and vae")
```

```
/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning:
  (prop.get_family(), self.defaultFamily[fonttext]))
```



The Laplacian variance increases with increased focus of an image or decreases with increased blur. Furthermore, images with a smaller amount of edges tend to have a smaller Laplacian variance (the Laplacian kernel is often used for edge detection in images). Therefore we first have to analyze the Laplacian variances for digit classes 0-9 in the MNIST test set before we can compare blur differences in generated images:

### 0.13 Model 1

- Train VAE
- Test Classifier

In [0]:

```
torch.manual_seed(5)
```

Out[0]: <torch.\_C.Generator at 0x7f535c0d2c10>

In [0]: `class Classifier(nn.Module):`

```
    def __init__(self):
```

```
        super(Classifier, self).__init__()
```

```
        ## Define NN
```

```
        self.fc1 = nn.Linear(10, 10)
```

```
    def forward(self, x):
```



```

    ## flat input features
    x = x.view(-1, self.num_flat_features(x))
    x = self.fc1(x)

    return F.log_softmax(x, dim=1)

def num_flat_features(self, x):
    size = x.size()[1:]
    # all dimensions except the batch dimension
    num_features = 1
    for s in size:
        num_features *= s
    return num_features

```

```

In [0]: ## Training
        from tqdm import trange

        criterion = nn.CrossEntropyLoss()
        classifier = Classifier()
        if is_cuda:
            classifier = classifier.to(device)

        # Loss and optimizer
        learning_rate = 0.001
        momentum = 0.9
        criterion = nn.CrossEntropyLoss()
        optimizer = torch.optim.Adam(classifier.parameters(),
                                      lr=learning_rate)
        scheduler = lr_scheduler.StepLR(optimizer,
                                         step_size = 7, gamma = 0.1)

        print(classifier)

```

```

Classifier(
  (fc1): Linear(in_features=10, out_features=10, bias=True)
)

```

```

In [0]: ## Train Classifier with pretrained vae
        vae_parameters = list(model.named_parameters())
        for name, param in vae_parameters:
            param.requires_grad = False

```

```

In [0]: def train_classifier_epoch(epoch):
        classifier.train()
        metric = AccumulatedAccuracyMetric()
        losses = RunningAverage()
        for idx, (data, labels) in enumerate(train_loader):
            data = data.to(device)
            labels = labels.to(device)

```

```

recon_batch, mu, logvar = model(data)
## classifier, pass latent vector
outputs = classifier(mu)
classifier_loss = criterion(outputs, labels)

optimizer.zero_grad()
classifier_loss.backward()
optimizer.step()

classifier_loss /= data.size(0)
losses.update(classifier_loss)

metric(outputs, labels)

return losses(), metric

```

```

In [0]: ## Test Epoch
        """
        Test, classifier on learnt features
        """
def test_classifier_epoch(epoch):
    classifier.eval()
    metric = AccumulatedAccuracyMetric()
    losses = RunningAverage()
    for idx, (data, labels) in enumerate(test_loader):
        data= data.to(device)
        labels = labels.to(device)

        recon_batch, mu, logvar = model(data)
        ## classifier, pass latent vector
        outputs = classifier(mu)
        classifier_loss = criterion(outputs, labels)

        classifier_loss /= data.size(0)
        losses.update(classifier_loss)

        metric(outputs, labels)

    return losses(), metric

```

```

In [0]: train_losses = []
        train_accuracy = []

```

```

test_losses = []
test_accuracy = []
n_epochs = 50
for epoch in range(1, n_epochs):

    # Train stage
    train_loss, metric = train_classifier_epoch(epoch)
    train_losses.append(train_loss)
    train_accuracy.append(metric.value())

    message = 'Epoch: {}/{}. Train set: Average loss: {:.4f}'
        .format(epoch + 1, n_epochs, train_loss)
    message += '\t Average Accuracy: \t{}: {}'.format(
        metric.name(), metric.value())
    print(message)

    val_loss, metrics = test_classifier_epoch(epoch)
    test_losses.append(val_loss)
    test_accuracy.append(metrics.value())

    message += '\nEpoch: {}/{}. Test set: Average loss: {:.4f}'
        .format(epoch + 1, n_epochs, val_loss)

    message += '\t Average Accuracy: \t{}: {}'.format(
        metrics.name(), metrics.value())

    print(message)

```

Epoch: 2/50. Train set: Average loss: 0.0134	Average Accuracy:	Accuracy: 52.31
Epoch: 2/50. Train set: Average loss: 0.0134	Average Accuracy:	Accuracy: 52.31
Epoch: 2/50. Test set: Average loss: 0.0099	Average Accuracy:	Accuracy: 85.35
Epoch: 3/50. Train set: Average loss: 0.0071	Average Accuracy:	Accuracy: 88.91
Epoch: 3/50. Train set: Average loss: 0.0071	Average Accuracy:	Accuracy: 88.91
Epoch: 3/50. Test set: Average loss: 0.0063	Average Accuracy:	Accuracy: 91.46
Epoch: 4/50. Train set: Average loss: 0.0048	Average Accuracy:	Accuracy: 91.64
Epoch: 4/50. Train set: Average loss: 0.0048	Average Accuracy:	Accuracy: 91.64
Epoch: 4/50. Test set: Average loss: 0.0045	Average Accuracy:	Accuracy: 92.62
Epoch: 5/50. Train set: Average loss: 0.0037	Average Accuracy:	Accuracy: 92.61
Epoch: 5/50. Train set: Average loss: 0.0037	Average Accuracy:	Accuracy: 92.61
Epoch: 5/50. Test set: Average loss: 0.0036	Average Accuracy:	Accuracy: 93.16
Epoch: 6/50. Train set: Average loss: 0.0031	Average Accuracy:	Accuracy: 93.11
Epoch: 6/50. Train set: Average loss: 0.0031	Average Accuracy:	Accuracy: 93.11
Epoch: 6/50. Test set: Average loss: 0.0032	Average Accuracy:	Accuracy: 93.48
Epoch: 7/50. Train set: Average loss: 0.0027	Average Accuracy:	Accuracy: 93.40
Epoch: 7/50. Train set: Average loss: 0.0027	Average Accuracy:	Accuracy: 93.40
Epoch: 7/50. Test set: Average loss: 0.0028	Average Accuracy:	Accuracy: 93.76
Epoch: 8/50. Train set: Average loss: 0.0024	Average Accuracy:	Accuracy: 93.62

[illegible]

[illegible]

Epoch: 40/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 40/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.73
Epoch: 41/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 41/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 41/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.73
Epoch: 42/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 42/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 42/50. Test set: Average loss: 0.0017	Average Accuracy:	Accuracy: 94.71
Epoch: 43/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 43/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 43/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.69
Epoch: 44/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 44/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 44/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.71
Epoch: 45/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 45/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 45/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.71
Epoch: 46/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 46/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 46/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 94.73
Epoch: 47/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 47/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 47/50. Test set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 48/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 48/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.73
Epoch: 48/50. Test set: Average loss: 0.0017	Average Accuracy:	Accuracy: 94.72
Epoch: 49/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 49/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.74
Epoch: 49/50. Test set: Average loss: 0.0017	Average Accuracy:	Accuracy: 94.7
Epoch: 50/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.76
Epoch: 50/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.76
Epoch: 50/50. Test set: Average loss: 0.0015	Average Accuracy:	Accuracy: 94.72

### 0.13.1 Save Classifier

```
In [0]: path =save_dir+"/classifier_model_state_"+str(epoch)+".pth"
        torch.save(classifier.state_dict(),path)


In [0]: path ="classifier_model_state_"+str(49)+".pth"
        load_model(classifier, path)
```

### 0.14 Plots

```
In [0]: plt.plot(range(1,n_epochs),train_accuracy)
        plt.plot(range(1,n_epochs),test_accuracy)
        plt.title("Classification accuracy of classifier")
        plt.legend(["train accuracy","test accuracy"])
```

Out[0]: <matplotlib.legend.Legend at 0x7fb7497064a8>


```
/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning  
(prop.get_family(), self.defaultFamily[fonttext]))
```



```
In [0]: plt.plot(range(1,n_epochs),train_losses)  
plt.plot(range(1,n_epochs),test_losses)  
plt.title("Classification loss of classifier")  
plt.legend(["train loss","test loss"])
```

Out[0]: <matplotlib.legend.Legend at 0x7fb748f34898>

```
/home/chivukula_manju/yes/lib/python3.6/site-packages/matplotlib/font_manager.py:1328: UserWarning  
(prop.get_family(), self.defaultFamily[fonttext]))
```



In [0]:

In [0]:

### 0.15 Model 3

```
In [0]: classifier = Classifier()
        model = ConvVAE(Params.nb_latents)

        if is_cuda:
            classifier = classifier.to(device)
            model = model.cuda()

        ## VAE is trained
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.Adam(list(model.parameters())
                                + list(classifier.parameters()), lr=1e-3)

        """
        Train VAE + but not classifier
        """
        def train_classifier_all(epoch):

            model.train()
            classifier.train()
```



```

metric = AccumulatedAccuracyMetric()
losses = RunningAverage()
for idx, (data, labels) in enumerate(train_loader):

    data = data.to(device)
    labels = labels.to(device)

    recon_batch, mu, logvar = model(data)
    kld, loss = loss_function(recon_batch.squeeze().view(-1,28*28),
                             data, mu, logvar)

    ## classifier, pass latent vector
    outputs = classifier(mu)
    classifier_loss = criterion(outputs, labels)

    ## Add all losses.
    loss = loss + classifier_loss

    ## parameter update
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()

    loss /=len(data)

    losses.update(loss)

    metric(outputs, labels)

return losses(), metric


## Test Epoch
"""
Test, classifier on learnt features
"""
def test_with_all_training(epoch):
    classifier.eval()
    model.eval()
    metric = AccumulatedAccuracyMetric()
    #losses = RunningAverage()
    for idx, (data, labels) in enumerate(test_loader):

        data, labels = data.to(device), labels.to(device)

        recon_batch, mu, logvar = model(data)

```

```

    ## classifier, pass latent vector
    outputs = classifier(mu)

```

```

    metric(outputs, labels)

```

```

    return 0.0, metric

```

```

In [0]: import warnings

```

```

warnings.filterwarnings("ignore")

```

```

train_losses = []

```

```

train_accuracy = []

```

```

test_losses = []

```

```

test_accuracy = []

```

```

n_epochs = 50

```

```

for epoch in range(1, n_epochs):

```

```

    # Train stage

```

```

    train_loss, metric = train_classifier_all(epoch)

```

```

    train_losses.append(train_loss)

```

```

    train_accuracy.append(metric.value())

```

```

    message = 'Epoch: {}/{}. Train set: Average loss: {:.4f}'
        .format(epoch + 1, n_epochs, train_loss)

```

```

    message += '\t Average Accuracy: \t{:}:'
        .format(metric.name(), metric.value())

```

```

    print(message)

```

```

    val_loss, metrics = test_with_all_training(epoch)

```

```

    test_losses.append(val_loss)

```

```

    test_accuracy.append(metrics.value())

```

```

    message += '\nEpoch: {}/{}. Test set: Average loss: {:.4f}'
        .format(epoch + 1, n_epochs, val_loss)

```

```

    message += '\t Average Accuracy: \t{:}:'
        .format(metrics.name(), metrics.value())

```

```

    print(message)

```

```

Epoch: 2/50. Train set: Average loss: -7492.7246

```

```

Epoch: 2/50. Train set: Average loss: -7492.7246

```

```

Epoch: 2/50. Test set: Average loss: 0.0000

```

```

Epoch: 3/50. Train set: Average loss: -8900.7314

```

```

Epoch: 3/50. Train set: Average loss: -8900.7314

```

```

Average Accuracy:

```

```

Average Accuracy:

```

```

Average Accuracy:

```

```

Average Accuracy:

```

```

Average Accuracy:

```

```

Accuracy: 2

```

```

Accuracy: 2

```

```

Accuracy: 52.34

```

```

Accuracy: 6

```

```

Accuracy: 6

```

Epoch: 3/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 75.07
Epoch: 4/50. Train set: Average loss: -9718.1816	Average Accuracy:	Accuracy: 8
Epoch: 4/50. Train set: Average loss: -9718.1816	Average Accuracy:	Accuracy: 8
Epoch: 4/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 84.7
Epoch: 5/50. Train set: Average loss: -10135.1260	Average Accuracy:	Accuracy: 8
Epoch: 5/50. Train set: Average loss: -10135.1260	Average Accuracy:	Accuracy: 8
Epoch: 5/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 88.17
Epoch: 6/50. Train set: Average loss: -10354.4287	Average Accuracy:	Accuracy: 8
Epoch: 6/50. Train set: Average loss: -10354.4287	Average Accuracy:	Accuracy: 8
Epoch: 6/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.86
Epoch: 7/50. Train set: Average loss: -10497.6748	Average Accuracy:	Accuracy: 8
Epoch: 7/50. Train set: Average loss: -10497.6748	Average Accuracy:	Accuracy: 8
Epoch: 7/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 91.05
Epoch: 8/50. Train set: Average loss: -10583.9287	Average Accuracy:	Accuracy: 9
Epoch: 8/50. Train set: Average loss: -10583.9287	Average Accuracy:	Accuracy: 9
Epoch: 8/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 90.33
Epoch: 9/50. Train set: Average loss: -10666.1738	Average Accuracy:	Accuracy: 9
Epoch: 9/50. Train set: Average loss: -10666.1738	Average Accuracy:	Accuracy: 9
Epoch: 9/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.25
Epoch: 10/50. Train set: Average loss: -10721.1650	Average Accuracy:	Accuracy:
Epoch: 10/50. Train set: Average loss: -10721.1650	Average Accuracy:	Accuracy:
Epoch: 10/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.61
Epoch: 11/50. Train set: Average loss: -10751.0410	Average Accuracy:	Accuracy:
Epoch: 11/50. Train set: Average loss: -10751.0410	Average Accuracy:	Accuracy:
Epoch: 11/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.88
Epoch: 12/50. Train set: Average loss: -10840.1602	Average Accuracy:	Accuracy:
Epoch: 12/50. Train set: Average loss: -10840.1602	Average Accuracy:	Accuracy:
Epoch: 12/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.58
Epoch: 13/50. Train set: Average loss: -10870.8945	Average Accuracy:	Accuracy:
Epoch: 13/50. Train set: Average loss: -10870.8945	Average Accuracy:	Accuracy:
Epoch: 13/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.55
Epoch: 14/50. Train set: Average loss: -10885.9775	Average Accuracy:	Accuracy:
Epoch: 14/50. Train set: Average loss: -10885.9775	Average Accuracy:	Accuracy:
Epoch: 14/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.73
Epoch: 15/50. Train set: Average loss: -10899.6182	Average Accuracy:	Accuracy:
Epoch: 15/50. Train set: Average loss: -10899.6182	Average Accuracy:	Accuracy:
Epoch: 15/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.86
Epoch: 16/50. Train set: Average loss: -10924.0127	Average Accuracy:	Accuracy:
Epoch: 16/50. Train set: Average loss: -10924.0127	Average Accuracy:	Accuracy:
Epoch: 16/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.96
Epoch: 17/50. Train set: Average loss: -10940.0918	Average Accuracy:	Accuracy:
Epoch: 17/50. Train set: Average loss: -10940.0918	Average Accuracy:	Accuracy:
Epoch: 17/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.11
Epoch: 18/50. Train set: Average loss: -10950.1260	Average Accuracy:	Accuracy:
Epoch: 18/50. Train set: Average loss: -10950.1260	Average Accuracy:	Accuracy:
Epoch: 18/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.2
Epoch: 19/50. Train set: Average loss: -10966.4785	Average Accuracy:	Accuracy:
Epoch: 19/50. Train set: Average loss: -10966.4785	Average Accuracy:	Accuracy:

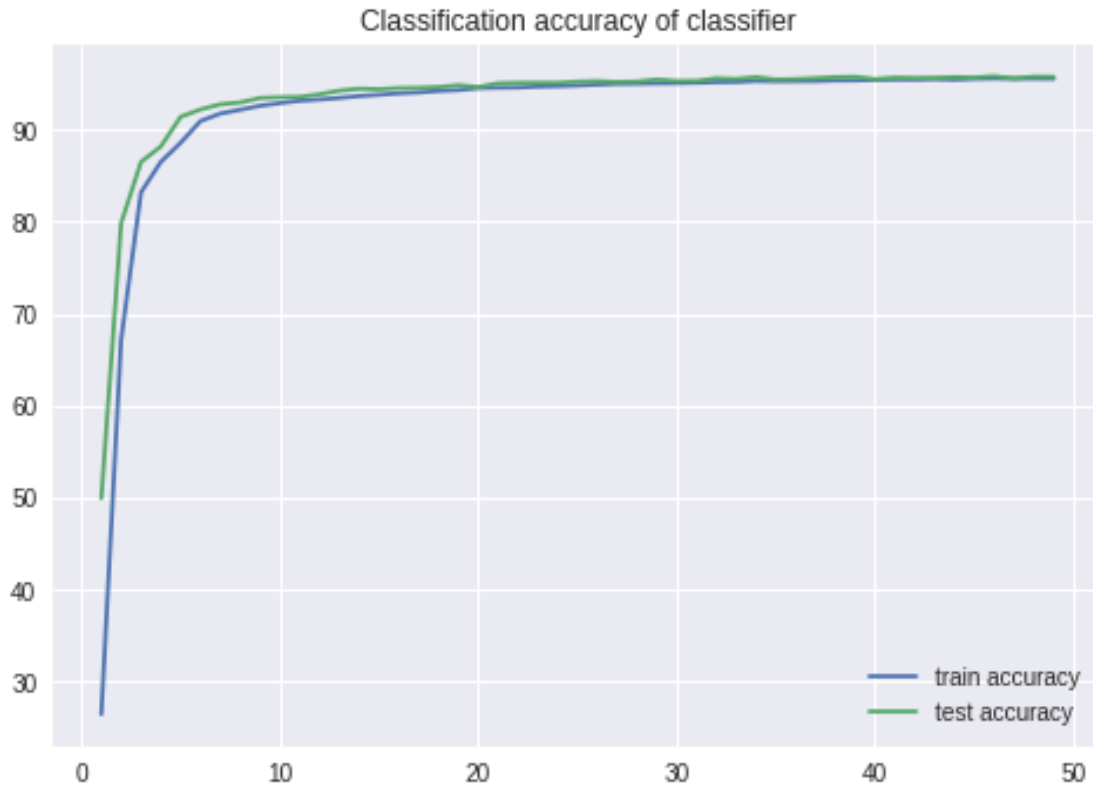
[illegible]

[illegible]

```
In [0]: plt.style.use("seaborn")
plt.plot(range(1,n_epochs),train_accuracy)
plt.plot(range(1,n_epochs),test_accuracy)
plt.title("Classification accuracy of classifier")
plt.legend(["train accuracy","test accuracy"])

print(max(train_accuracy),max(test_accuracy))
```

Out[0]: <matplotlib.legend.Legend at 0x7f996178d4a8>



```
In [0]: path = "tsne_bvae+classifier+vae+train.png"
def visualize_tsne(X, labels, model, path):
    # Compute latent space representation
    print("Computing latent space projection...")

    X_encoded, _ = model.encode(X)

    # Compute t-SNE embedding of latent space
    tsne = manifold.TSNE(n_components=2, init='pca',
                        random_state=0)
    X_tsne = tsne.fit_transform(X_encoded.data.detach().cpu())
```

```


# Plot images according to t-sne embedding
fig, ax = plt.subplots()

plt.scatter(X_tsne[:,0], X_tsne[:,1], c=labels,
            cmap=plt.cm.get_cmap("viridis", 10))
plt.colorbar(ticks=range(10))
fig.savefig(path, dpi=fig.dpi)

visualize_tsne(test_data[:5000].to(device), test_labels[:5000], model, path)

```

Computing latent space projection...



In [0]:

In [0]: *## latent traversals*

```

testpoint1, testpointlabel1 = train_loader.dataset[0]
testpoint2, testpointlabel2 = train_loader.dataset[1]
testpoint3, testpointlabel3 = train_loader.dataset[2]
filename = 'beta-results/vae_traversal_testpoint_5.png'

```

```
traverse_latents(model, testpoint1, Params.nb_latents, filename)

filename = 'beta-results/vae_traversal_testpoint_0.png'
traverse_latents(model, testpoint2, Params.nb_latents, filename)

filename = 'beta-results/vae_traversal_testpoint_4.png'
traverse_latents(model, testpoint3, Params.nb_latents, filename)
```

In [0]: