

Plots

April 25, 2019

```
In [0]: from glob import glob
import os
import numpy as np
import matplotlib.pyplot as plt
import shutil
from torchvision import transforms
from torchvision import models
import torchvision
import torch
from torch.autograd import Variable
import torch.nn as nn
from torch.optim import lr_scheduler
from torch import optim
from torchvision.datasets import ImageFolder
from torchvision.utils import make_grid
import time
from torchvision import datasets, transforms
import torch.nn.functional as F
import torch.optim as optim

import os
from torchvision.utils import save_image

## Plotting library
from matplotlib.offsetbox import OffsetImage, AnnotationBbox
from matplotlib.backends.backend_agg import FigureCanvasAgg as FigureCanvas

from scipy.stats import norm
from sklearn import manifold

%matplotlib inline
import seaborn as sns
plt.style.use("seaborn")

print('Torch', torch.__version__, 'CUDA', torch.version.cuda)
print('Device:', torch.device('cuda:0'))
print(torch.cuda.is_available())
```

```

is_cuda = torch.cuda.is_available()
device = torch.device ( "cuda:0" if torch.cuda.is_available () else "cpu" )

```

Torch 1.0.1.post2 CUDA 9.0.176

Device: cuda:0

True

In [0]: *## Data loader*

```

def get_dataloaders(path, batch_size, pad = False):
    path = os.path.join(path, "MNIST/data")
    print(path)
    transform = transforms.Compose([transforms.ToTensor(),
                                     transforms.Normalize((0.1307,), (0.3081,))])
    train_dataset = datasets.MNIST(root = path,
                                    train = True, download = False,
                                    transform = transform)
    train_loader = torch.utils.data.DataLoader(dataset = train_dataset,
                                                batch_size = batch_size,
                                                shuffle = True)

    test_dataset = datasets.MNIST(root=path,train=False,
                                   transform=transform,download=True)
    test_loader = torch.utils.data.DataLoader(test_dataset,
                                                batch_size=batch_size,shuffle=True)

    return (train_loader, test_loader)

```

In [0]: *"""*

Load model
"""

```

def load_model(model, filepath):
    state = torch.load(filepath)
    model.load_state_dict(state['model'])
    model.cuda()

```

```

def load_model_state(model, filepath):
    state = torch.load(filepath)
    model.load_state_dict(state)
    model.cuda()

```

In [0]:

"""

Traverse Latents
"""

```

def traverse_latents(model, datapoint, nb_latents, file):

```

```

model.eval()
datapoint = datapoint.to(device)
if isinstance(model, ConvVAE):
    datapoint = datapoint.unsqueeze(0)
    mu, _ = model.encode(datapoint)
else:
    datapoint = datapoint.unsqueeze(0)
    mu, _ = model.encoder(datapoint)

recons = torch.zeros((7, nb_latents, 28, 28))

for zi in range(nb_latents):
    muc = mu.squeeze().clone()
    for i, val in enumerate(np.linspace(-3, 3, 7)):
        muc[zi] = val
        recon = model.decode(muc).cpu()
        if isinstance(model, ConvVAE) else model.decoder(muc).cpu()
        recons[i, zi] = recon.view(28, 28)

torchvision.utils.save_image(recons.view(-1, 1, 28, 28),
                             file, nrow=nb_latents, pad_value=1)

import cv2
import numpy

def laplacian_variance(images):
    return [cv2.Laplacian(image.numpy(), cv2.CV_32F).var() for image in images]

def laplacian_variance_numpy(images):
    return [cv2.Laplacian(image, cv2.CV_32F).var() for image in images]

In [0]: class ConvVAE(nn.Module):
        def __init__(self, nb_latents):
            super(ConvVAE, self).__init__()
            self.conv1 = nn.Sequential(
                nn.Conv2d(1, 32,
                           kernel_size=5, stride=1, padding=2),
                nn.ReLU(),
                nn.MaxPool2d(kernel_size=2, stride=2))
            self.conv2 = nn.Sequential(
                nn.Conv2d(32, 64,
                           kernel_size=5, stride=1, padding=2),
                nn.ReLU(),
                nn.MaxPool2d(kernel_size=2, stride=2))

```

```

self.conv3 = nn.Sequential(
    nn.Conv2d(64, 128,
              kernel_size=5, stride=1, padding=2),
    nn.ReLU(),
    nn.MaxPool2d(kernel_size=2, stride=2))

self.conv4 = nn.Sequential(
    nn.Conv2d(128, 256,
              kernel_size=2, stride=1),
    nn.ReLU())

self.fc1 = nn.Linear(1024, 256)

self.fc_mean = nn.Linear(256, nb_latents)
self.fc_std = nn.Linear(256, nb_latents)

self.fc2 = nn.Linear(nb_latents, 256)
self.fc3 = nn.Linear(256, 1024)

self.fc4 = nn.Linear(1024, 7*7*64)

self.deconv1 = nn.ConvTranspose2d(64, 32,
                                   kernel_size=2, stride=2)
self.deconv2 = nn.ConvTranspose2d(32, 1,
                                   kernel_size=2, stride=2)

self.relu = nn.ReLU()
self.sigmoid = nn.Sigmoid()

def encode(self, x):
    x = (self.conv1(x))
    x = (self.conv2(x))
    x = (self.conv3(x))
    x = (self.conv4(x))
    x = x.reshape(x.size(0), -1)
    x = self.relu(self.fc1(x))
    return self.fc_mean(x), self.fc_std(x)

def reparameterize(self, mu, logvar):
    if self.training:
        std = logvar.mul(0.5).exp_()
        eps = Variable(std.data.new(std.size()).normal_())
        return eps.mul(std).add_(mu)
    else:
        return mu

```

```

def decode(self, z):
    x = self.relu(self.fc2(z))
    x = self.relu(self.fc3(x))
    x = self.relu(self.fc4(x))
    x = self.relu(self.deconv1(x.view(-1, 64, 7, 7)))
    x = self.deconv2(x)
    return self.sigmoid(x)

def forward(self, x):
    mu, logvar = self.encode(x)
    z = self.reparameterize(mu, logvar)
    return self.decode(z), mu, logvar

In [0]: vae_model = ConvVAE(10)
        if is_cuda:
            vae_model = vae_model.to(device)

        ## load model
        checkpoint_dir = './vae-final/checkpoints'
        filename = checkpoint_dir + "/model_state_ "+str(99)+ ".pth"
        load_model(vae_model, filename)

In [0]: class Params:
        nb_latents = 10
        batch_size = 128
        epochs = 100
        log_interval = 100
        save_interval = 1000
        data_dir = "./vae-classifier/"
        train_loader, test_loader = get_dataloaders(Params.data_dir, Params.batch_size)

        ./vae-classifier/MNIST/data

In [0]: ## TSNE for VAE
        ## Test TSNE plot for reconstruction on 1000 test samples

        transformation = transforms.Compose([transforms.ToTensor(),
                                             transforms.Normalize((0.1307,), (0.3081,))])
        train_dataset = datasets.MNIST(os.path.join(Params.data_dir, "MNIST/data"),
                                       train=True, transform=transformation, download=False)
        testing_tsne = torch.utils.data.DataLoader(train_dataset,
                                                    batch_size=len(train_dataset), shuffle=True)
        test_data, test_labels = next(iter(testing_tsne))[:10000]

In [0]: path = "./vae-final/samples/"+'/latent_space.png'
        def visualize_tsne(X, labels, model, path):
            # Compute latent space representation

```

```

print("Computing latent space projection...")

#if isinstance(model,VAE_GAN_Generator):
    # X_encoded, _ = model.encoder(X)
#else:
X_encoded, _ = model.encode(X)


# Compute t-SNE embedding of latent space
tsne = manifold.TSNE(n_components=2, init='pca', random_state=0)
X_tsne = tsne.fit_transform(X_encoded.data.detach().cpu())

# Plot images according to t-sne embedding
fig, ax = plt.subplots()
plt.scatter(X_tsne[:,0], X_tsne[:,1],
            c=labels, cmap=plt.cm.get_cmap("viridis", 10))
plt.colorbar(ticks=range(10))
fig.savefig(path, dpi=fig.dpi)

visualize_tsne(test_data[:5000].to(device),test_labels[:5000],vae_model,path)

```

Computing latent space projection...




0.1 Beta VAE

```
In [0]: beta_vae = ConvVAE(Params.nb_latents)
        if is_cuda:
            beta_vae = beta_vae.to(device)

        ## load model
        #model_state_99.pth
        checkpoint_dir = 'beta-results'
        filename = checkpoint_dir + "/model_state_" + str(99) + ".pth"
        load_model_state(beta_vae, filename)

In [0]: path = "beta-results/latent-space.png"
        visualize_tsne(test_data[:5000].to(device),
                        test_labels[:5000], beta_vae, path)
```

Computing latent space projection...



In [0]:

0.2 VAE+GAN

```
In [0]: ## Load GAN Model
def load_Vae_gan(epoch, encoder, decoder, D):
    save_dir = "vae-gan"
    # restore models
    decoder.load_state_dict(torch.load(save_dir+
        '/VAE_GAN_decoder_%d.pth' % epoch))
    decoder.cuda()
    encoder.load_state_dict(torch.load(save_dir +
        '/VAE_GAN_encoder_%d.pth' % epoch))
    encoder.cuda()
    D.load_state_dict(torch.load(save_dir+
        '/VAE_GAN_D_%d.pth' % epoch))
    D.cuda()

In [0]: class Encoder(nn.Module):
    def __init__(self, input_channels, output_channels, representation_size = 32):
        super(Encoder, self).__init__()
        # input parameters
        self.input_channels = input_channels
        self.output_channels = output_channels #Params.nb_latents

        self.features = nn.Sequential(
            # nc x 32 x 32
            nn.Conv2d(self.input_channels,
                representation_size, 5, stride=2, padding=2),
            nn.BatchNorm2d(representation_size),
            nn.ReLU(),
            # hidden_size x 16 x 16
            nn.Conv2d(representation_size,
                representation_size*2, 5, stride=2, padding=2),
            nn.BatchNorm2d(representation_size * 2),
            nn.ReLU(),
            # hidden_size*2 x 8 x 8
            nn.Conv2d(representation_size*2,
                representation_size*4, 5, stride=2, padding=2),
            nn.BatchNorm2d(representation_size * 4),
            nn.ReLU())
            # hidden_size*4 x 4 x 4

        self.mean = nn.Sequential(
            nn.Linear(representation_size*4*4*4, 1024),

            nn.ReLU(),
            nn.Linear(1024, output_channels))

        self.logvar = nn.Sequential(
            nn.Linear(representation_size*4*4*4, 1024),
```



```

        nn.ReLU(),
        nn.Linear(1024, output_channels))

def forward(self, x):
    batch_size = x.size()[0]

    hidden_representation = self.features(x)
    hidden_representation = hidden_representation
        .view(-1, self.num_flat_features(hidden_representation))

    mean = self.mean(hidden_representation)
    logvar = self.logvar(hidden_representation)

    return mean, logvar

def hidden_layer(self, x):
    batch_size = x.size()[0]
    output = self.features(x)
    return output

def num_flat_features(self, x):
    size = x.size()[1:] # all dimensions except the batch dimension
    num_features = 1
    for s in size:
        num_features *= s
    return num_features

```

```

In [0]: class Decoder(nn.Module):
    def __init__(self, input_size):
        super(Decoder, self).__init__()

        self.input_size = input_size ##

        self.fc1 = nn.Linear(self.input_size, 256)
        self.fc2 = nn.Linear(256, 1024)
        self.fc3 = nn.Linear(1024, 7*7*64)

        self.deconv1 = nn.ConvTranspose2d(64, 32, kernel_size=2, stride=2)
        self.deconv2 = nn.ConvTranspose2d(32, 1, kernel_size=2, stride=2)

        # 1 x 28 x 28
        self.activation = nn.Sigmoid()
        self.relu = nn.Tanh()

    def forward(self, z):

```

```

bs = z.size()[0]

x = self.relu(self.fc1(z))
x = self.relu(self.fc2(x))
x = self.relu(self.fc3(x))
x = self.relu(self.deconv1(x.view(-1, 64, 7, 7)))
x = self.deconv2(x)

return self.activation(x)

```

```

In [0]: class VAE_GAN_Generator(nn.Module):
    def __init__(self, input_channels, hidden_size):
        super(VAE_GAN_Generator, self).__init__()
        self.input_channels = input_channels
        self.hidden_size = hidden_size

        self.encoder = Encoder(input_channels, hidden_size)
        self.decoder = Decoder(hidden_size)

    def forward(self, x):
        batch_size = x.size()[0]
        mean, logvar = self.encoder(x)
        std = logvar.mul(0.5).exp_()

        reparametrized_noise = Variable(torch.randn(
            (batch_size, self.hidden_size))).cuda()

        reparametrized_noise = mean + std * reparametrized_noise

        rec_images = self.decoder(reparametrized_noise)

        return mean, logvar, rec_images

```

```

In [0]: class Discriminator(nn.Module):
    def __init__(self, input_channels, representation_size = 32):
        super(Discriminator, self).__init__()
        self.input_channels = input_channels
        dim = 128 * 4 * 4
        self.main = nn.Sequential(
            # nc x 32 x 32
            nn.Conv2d(self.input_channels,
                representation_size, 5, stride=2, padding=2),
            nn.BatchNorm2d(representation_size),
            nn.LeakyReLU(0.2),
            # hidden_size x 16 x 16
            nn.Conv2d(representation_size,
                representation_size*2, 5, stride=2, padding=2),
            nn.BatchNorm2d(representation_size * 2),

```

```

        nn.LeakyReLU(0.2),
        # hidden_size*2 x 8 x 8
        nn.Conv2d(representation_size*2,
                  representation_size*4, 5, stride=2, padding=2),
        nn.BatchNorm2d(representation_size * 4),
        nn.LeakyReLU(0.2))

    self.lth_features = nn.Sequential(
        nn.Linear(dim, 2048),
        nn.LeakyReLU(0.2))

    self.sigmoid_output =
    nn.Sequential(nn.Linear(2048,1),nn.Sigmoid())

    def forward(self, x):
        batch_size = x.size()[0]
        features = self.main(x)
        lth_rep = self.
        lth_features(features.view(batch_size, -1))
        output = self.sigmoid_output(lth_rep)
        return output

    def similarity(self, x):
        batch_size = x.size()[0]
        features = self.main(x)
        lth_rep = self.
        lth_features(features.view(batch_size, -1))
        return lth_rep

```

In [0]: # define constant

```

input_channels = 1
hidden_size = 10
max_epochs = 250
lr = 3e-4

```

```

beta = 5
alpha = 0.1
gamma = 15

```

In [0]: G = VAE_GAN_Generator(input_channels, hidden_size).cuda()
D = Discriminator(input_channels).cuda()

```

criterion = nn.BCELoss()
criterion.cuda()

```

```

opt_enc = optim.RMSprop(G.encoder.parameters(), lr=lr)
opt_dec = optim.RMSprop(G.decoder.parameters(), lr=lr)
opt_dis = optim.RMSprop(D.parameters(), lr=lr * alpha)

```

```
In [0]: load_Vae_gan(210, G.encoder, G.decoder, D)

In [ ]: path = "vae-gan/latent-space.png"
        visualize_tsne(test_data[:5000].to(device), test_labels[:5000], G, path)
```

0.3 PLOTS

```
In [0]: ## TSNE Plot of VAE, Beta VAE and VAE + GAN
import matplotlib.pyplot as plt
from IPython.display import Image
import matplotlib.image as mpimg
import matplotlib.gridspec as gridspec

rows = 2
columns = 3

fig_height = 5.
length_x_axis = 100
length_y_axis = 45
height = length_y_axis * rows
width = length_x_axis * columns

plot_aspect_ratio=2

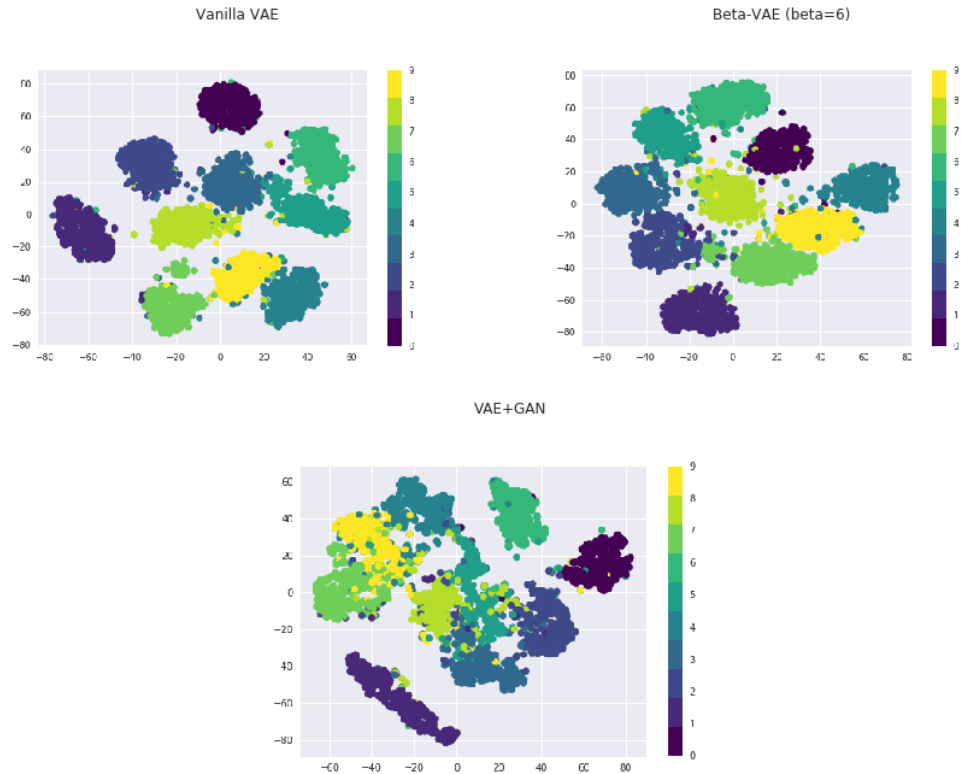
fig = plt.figure(figsize = (16,12))
gs = gridspec.GridSpec(2,2)
gs.update(wspace=0.025, hspace=0.05)

## Train
ax1 = plt.subplot(gs[0, 0])
vae_tsne = "vae-final/samples" + '/latent_space.png'
img=mpimg.imread(vae_tsne)
imgplot = plt.imshow(img)
plt.axis("off")
plt.title("Vanilla VAE")

ax2 = plt.subplot(gs[0, 1])
betavae_tsne = "beta-results/latent-space.png"
img=mpimg.imread(betavae_tsne)
imgplot = plt.imshow(img)
plt.axis("off")
plt.title("Beta-VAE (beta=6)")

ax3 = plt.subplot(gs[1,0:])
vae_gan_tsne = "./vae-gan/latent-space.png"
img=mpimg.imread(vae_gan_tsne)
imgplot = plt.imshow(img)
plt.title("VAE+GAN")
```

```
plt.axis("off")
plt.savefig("vae-final/tsne-results.png")
```



0.4 Reconstructions

```
In [0]: def reconstruction(model, data, path):
        n = min(data.size(0), 10)
        recon_batch, mu, logvar = model(data)
        comparison = torch.cat([data[:n], recon_batch[:n]])
        save_image(comparison.cpu(), path, nrow=Params.nb_latents, pad_value=1)
```

```
train_batch, train_label = next(iter(train_loader))
test_batch, test_label = next(iter(test_loader))
```

```
In [0]: filename = os.path.join("vae-final/", 'vae_reconstruction.png')
        reconstruction(vae_model, train_batch.to(device), filename)
```

```
In [0]: filename = os.path.join("vae-final/", 'beta_vae_reconstruction.png')
        reconstruction(beta_vae, train_batch.to(device), filename)
```

```

In [0]: def reconstruction_vae_gan(model, data, path):
        n = min(data.size(0), 10)
        _, _, rec_imgs = model(data)
        comparison = torch.cat([data[:n], rec_imgs[:n]])
        save_image(comparison.cpu(), path,
                    nrow=Params.nb_latents, pad_value=1)

        filename = os.path.join("vae-final/",
                                'vae_gan_reconstruction.png')
        reconstruction_vae_gan(G, train_batch.to(device), filename)

In [0]: rows = 2
        columns = 3

        fig_height = 5.
        length_x_axis = 100
        length_y_axis = 45
        height = length_y_axis * rows
        width = length_x_axis * columns

        plot_aspect_ratio = float(width)/float(height)

        fig = plt.figure(figsize=(10,3))

        gs = gridspec.GridSpec(2,2)
        gs.update(wspace = 0.23, hspace = 0.05)
        ## Train
        ax1 = plt.subplot(gs[0, 0])
        vae_tsne = "vae-final/vae_reconstruction.png"
        img=mpimg.imread(vae_tsne)
        imgplot = plt.imshow(img)
        plt.title("Vanilla VAE")
        plt.axis("off")
        ax2 = plt.subplot(gs[0, 1])
        betavae_tsne = "vae-final/beta_vae_reconstruction.png"
        img=mpimg.imread(betavae_tsne)
        imgplot = plt.imshow(img)
        plt.title("Beta-VAE (beta=6)")
        plt.axis("off")
        ax3 = plt.subplot(gs[1,:])
        vae_gan_tsne = "vae-final/vae_gan_reconstruction.png"
        img=mpimg.imread(vae_gan_tsne)
        imgplot = plt.imshow(img)
        plt.title("VAE+GAN")
        plt.axis("off")
        plt.savefig("vae-final/reconstruction_results.png")

```



0.5 Samples

```
In [0]: sample = torch.randn(64, Params.nb_latents).to(device)
```

```
In [0]: # Generate samples of reconstruction
vae_model.eval()
with torch.no_grad():
    sample_vae = vae_model.decode(sample).cpu()
    save_image(sample_vae.view(64, 1, 28, 28),
                'vae-final/vae_sample.png')
```

```
In [0]: # Generate samples of reconstruction
beta_vae.eval()
with torch.no_grad():
    sample_bvae = beta_vae.decode(sample).cpu()
    save_image(sample_bvae.view(64, 1, 28, 28),
                'vae-final/bvae_sample.png')
```

```
In [0]: # Generate samples of reconstruction
G.eval()
with torch.no_grad():
    sample_vaeg = G.decoder(sample).cpu()
    save_image(sample_vaeg.view(64, 1, 28, 28),
                'vae-final/vae_gan_sample.png')
```

```
In [0]: rows = 2
        columns = 3

        fig_height = 5.
        length_x_axis = 100
        length_y_axis = 45
        height = length_y_axis * rows
        width = length_x_axis * columns
```


```

plot_aspect_ratio= float(width)/float(height)

fig = plt.figure(figsize=(8,10))

gs = gridspec.GridSpec(2,2)
gs.update(wspace = 0.025, hspace = 0.0)
## Train
ax1 = plt.subplot(gs[0, 0])
vae_tsne = "vae-final/vae_sample.png"
img=mpimg.imread(vae_tsne)
imgplot = plt.imshow(img)
plt.title("Vanilla VAE")
plt.axis("off")
ax2 = plt.subplot(gs[0, 1])
betavae_tsne = "vae-final/bvae_sample.png"
img=mpimg.imread(betavae_tsne)
imgplot = plt.imshow(img)
plt.title("Beta-VAE (beta=6)")
plt.axis("off")
ax3 = plt.subplot(gs[1, :])
vae_gan_tsne = "vae-final/vae_gan_sample.png"
img=mpimg.imread(vae_gan_tsne)
imgplot = plt.imshow(img)
plt.title("VAE+GAN")
plt.axis("off")
plt.savefig("vae-final/sample_results.png")

```

0.6 Latent Traversal

```
In [0]: testpoint1, testpointlabel1 = train_loader.dataset[0]
        testpoint2, testpoint2label2 = train_loader.dataset[1]
        testpoint3, testpoint3label3 = train_loader.dataset[2]

        print(testpointlabel1, testpoint2label2, testpoint3label3)
```

5 0 4

```
In [0]: filename = 'vae-final/vae_traversal_testpoint_5.png'
        traverse_latents(vae_model, testpoint1, Params.nb_latents, filename)
        filename = 'vae-final/beta_vae_traversal_testpoint_5.png'
        traverse_latents(beta_vae, testpoint1, Params.nb_latents, filename)
```

```
In [0]: filename = 'vae-final/vae_gan_traversal_testpoint_5.png'
        traverse_latents(G, testpoint1, Params.nb_latents, filename)
```

```
In [0]: def display_three_plots(filename1, filename2, filename3, save_name):
```

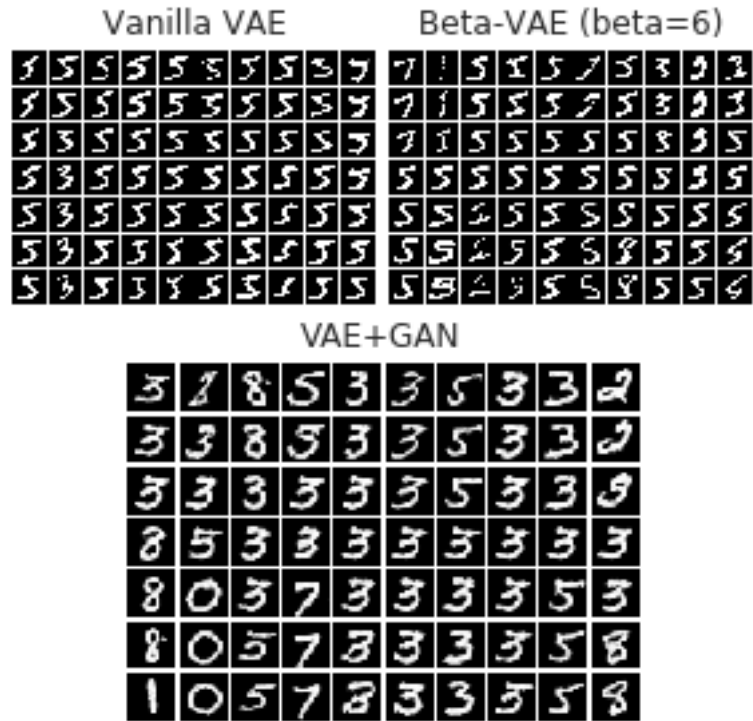
```
    fig = plt.figure(figsize=(5,5))
    gs = gridspec.GridSpec(2,2)
    gs.update(wspace = 0.025, hspace = 0.0)
```

```
    ## Train
```

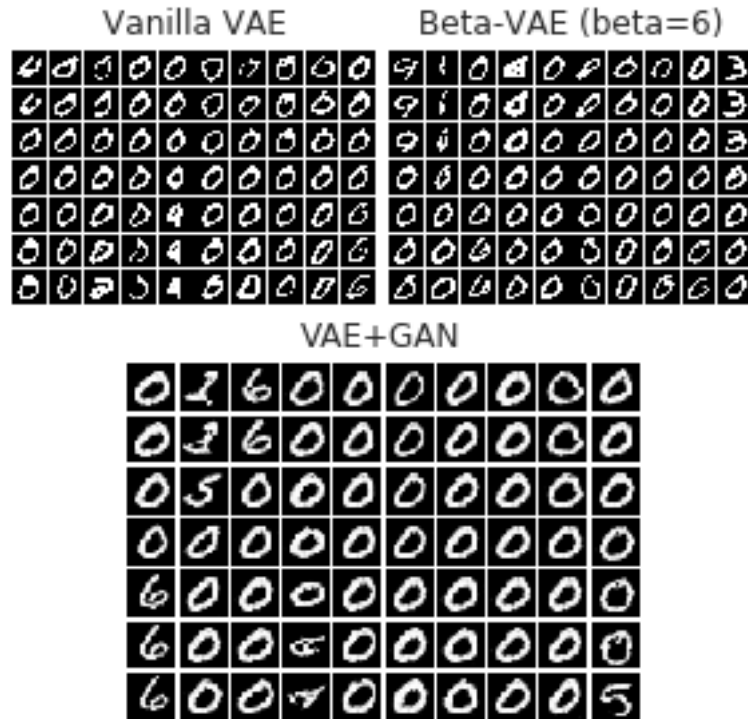
```
    ax1 = plt.subplot(gs[0, 0])
    vae_tsne = filename1
    img=mpimg.imread(vae_tsne)
    imgplot = plt.imshow(img)
    plt.title("Vanilla VAE")
    plt.axis("off")
    ax2 = plt.subplot(gs[0, 1])
    betavae_tsne = filename2
    img=mpimg.imread(betavae_tsne)
    imgplot = plt.imshow(img)
    plt.title("Beta-VAE (beta=6)")
    plt.axis("off")
    ax3 = plt.subplot(gs[1, :])
    vae_gan_tsne = filename3
    img=mpimg.imread(vae_gan_tsne)
    imgplot = plt.imshow(img)
    plt.title("VAE+GAN")
    plt.axis("off")
    plt.savefig("vae-final/"+save_name)
```

```
In [0]: filename1 = 'vae-final/vae_traversal_testpoint_5.png'
        filename2 = 'vae-final/beta_vae_traversal_testpoint_5.png'
        filename3 = 'vae-final/vae_gan_traversal_testpoint_5.png'
```


```
display_three_plots(filename1, filename2, filename3, "traversal.png")
```



```
In [0]: filename1 = 'vae-final/vae_traversal_testpoint_0.png'
        traverse_latents(vae_model, testpoint2, Params.nb_latents, filename1)
        filename2 = 'vae-final/beta_vae_traversal_testpoint_0.png'
        #traverse_latents(beta_vae, testpoint2, Params.nb_latents, filename2)
        filename3 = 'vae-final/vae_gan_traversal_testpoint_0.png'
        traverse_latents(G, testpoint2, Params.nb_latents, filename3)
        display_three_plots(filename1, filename2, filename3, "traversal0.png")
```



```
In [0]: filename1 = 'vae-final/vae_traversal_testpoint_4.png'
        traverse_latents(vae_model, testpoint3, Params.nb_latents, filename1)
        filename2 = 'vae-final/beta_vae_traversal_testpoint_4.png'
        #traverse_latents(beta_vae, testpoint3, Params.nb_latents, filename2)
        filename3 = 'vae-final/vae_gan_traversal_testpoint_4.png'
        traverse_latents(G, testpoint3, Params.nb_latents, filename3)
        display_three_plots(filename1, filename2, filename3, "traversal4.png")
```



```
In [0]: rows = 3
        columns = 3

        fig_height = 5.
        length_x_axis = 100
        length_y_axis = 100
        height = length_y_axis * rows
        width = length_x_axis * columns

        plot_aspect_ratio= 3

        fig = plt.figure(figsize=(fig_height * plot_aspect_ratio, fig_height ))
        gs = gridspec.GridSpec(rows,columns)

        def add_subplot(row, col, filename,title):
            ## Train
            ax1 = plt.subplot(gs[row, col])
            vae_tsne = filename
            img=mpimg.imread(vae_tsne)
            imgplot = plt.imshow(img)
            plt.title(title)
            plt.axis("off")
```

```

filename1 = 'vae-final/vae_traversal_testpoint_5.png'
filename2 = 'vae-final/beta_vae_traversal_testpoint_5.png'
filename3 = 'vae-final/vae_gan_traversal_testpoint_5.png'

filename4 = 'vae-final/vae_traversal_testpoint_0.png'
filename5 = 'vae-final/beta_vae_traversal_testpoint_0.png'
filename6 = 'vae-final/vae_gan_traversal_testpoint_0.png'

filename7 = 'vae-final/vae_traversal_testpoint_4.png'
filename8 = 'vae-final/beta_vae_traversal_testpoint_4.png'
filename9 = 'vae-final/vae_gan_traversal_testpoint_4.png'

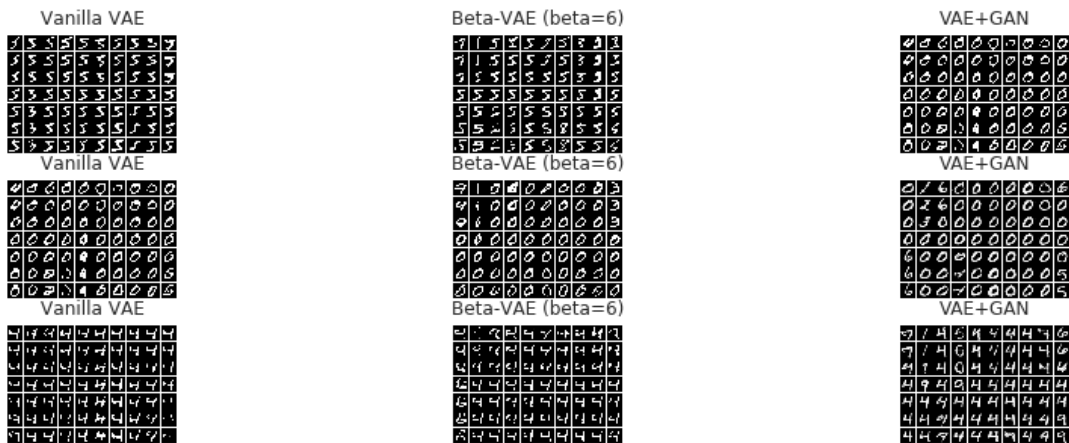
add_subplot(0,0,filename1, "Vanilla VAE")
add_subplot(0,1,filename2, "Beta-VAE (beta=6)")
add_subplot(0,2,filename3, "VAE+GAN")

add_subplot(1,0,filename4, "Vanilla VAE")
add_subplot(1,1,filename5, "Beta-VAE (beta=6)")
add_subplot(1,2,filename6, "VAE+GAN")

add_subplot(2,0,filename7, "Vanilla VAE")
add_subplot(2,1,filename8, "Beta-VAE (beta=6)")
add_subplot(2,2,filename9, "VAE+GAN")

plt.savefig("vae-final/all_traversal.png")

```



0.7 Laplacian Variance

In [0]:

```

In [0]: ## Import CV
import cv2
import numpy

def laplacian_variance(images):
    return [cv2.Laplacian(image.numpy(), cv2.CV_32F).var() for image in images]

def laplacian_variance_numpy(images):
    return [cv2.Laplacian(image, cv2.CV_32F).var() for image in images]

In [0]: x_test, y_test = test_data[:2000], test_labels[:2000]
        laplacian_variances = [laplacian_variance(x_test[y_test == i]) for i in range(10)]

In [0]: plt.boxplot(laplacian_variances, labels=range(10));
        plt.xlabel('Digit class')
        plt.ylabel('Laplacian variance')
        plt.title('Laplacian variance statistics per digit class');

```



```

In [0]: ## Vanilla VAE
def plot_laplacian_variances(lvs_1, lvs_2, label2, title, path):

```

```

plt.hist(lvs_1, bins=50, alpha=0.5 , label='Original images');
plt.hist(lvs_2, bins=50, alpha= 0.5 , label=label2);
plt.xlabel('Laplacian variance')
plt.title(title)
plt.legend()
plt.savefig(path)


```

```

In [0]: not_ones = y_test != 1
x_test_not_ones, y_test_not_ones = x_test[not_ones], y_test[not_ones]
with torch.no_grad():
    vae_reconstructions = np.empty(shape=(len(x_test_not_ones),1,28,28))
    indx = 0
    for i, (x,y) in enumerate(zip(x_test_not_ones,y_test_not_ones)):
        if y != 1:
            recon_batch, mu, logvar = vae_model(x.unsqueeze(1).to(device))
            vae_reconstructions[indx]=(recon_batch.squeeze(0).detach().cpu())
            indx+=1

not_ones = y_test != 1
lvs_1 = laplacian_variance(x_test[not_ones])
lvs_2 = laplacian_variance_numpy(np.array(vae_reconstructions, dtype=np.uint8))
filename1 = "vae-final/vae_laplacian_variance_reconstructed.png"
plot_laplacian_variances(lvs_1, lvs_2, "Images of VAE", "", filename1)

```




```

In [0]: not_ones = y_test != 1
        x_test_not_ones, y_test_not_ones = x_test[not_ones], y_test[not_ones]
        with torch.no_grad():
            bvae_reconstructions = np.empty(shape=(len(x_test_not_ones),1,28,28))
            indx = 0
            for i, (x,y) in enumerate(zip(x_test_not_ones,y_test_not_ones)):
                if y != 1:
                    recon_batch, mu, logvar = beta_vae(x.unsqueeze(1).to(device))
                    bvae_reconstructions[indx]
                    =(recon_batch.squeeze(0).detach().cpu())
                    indx+=1

        not_ones = y_test != 1
        lvs_1 = laplacian_variance(x_test[not_ones])
        lvs_2 = laplacian_variance_numpy(
            np.array(bvae_reconstructions, dtype=np.uint8))
        filename2 = "vae-final/bvae_laplacian_variance_reconstructed.png"
        plot_laplacian_variances(lvs_1, lvs_2,
            "Images of Beta-VAE (beta=6)", "", filename2)

```




```

In [0]: not_ones = y_test != 1
        x_test_not_ones, y_test_not_ones = x_test[not_ones], y_test[not_ones]
        with torch.no_grad():
            vae_gan_reconstructions =
            np.empty(shape=(len(x_test_not_ones),1,28,28))
            indx = 0
            for i, (x,y) in enumerate(zip(x_test_not_ones,y_test_not_ones)):
                if y != 1:
                    _, _, recon_batch = G(x.unsqueeze(1).to(device))
                    vae_gan_reconstructions[indx]
                    =(recon_batch.squeeze(0).detach().cpu())
                    indx+=1

        not_ones = y_test != 1
        lvs_1 = laplacian_variance(x_test[not_ones])
        lvs_2 = laplacian_variance_numpy(np.array(vae_gan_reconstructions, dtype=np.uint8))
        filename3 = "vae-final/vae_gan_laplacian_variance_reconstructed.png"
        plot_laplacian_variances(lvs_1, lvs_2,
                                "Images of VAE-GAN", "", filename3)

```



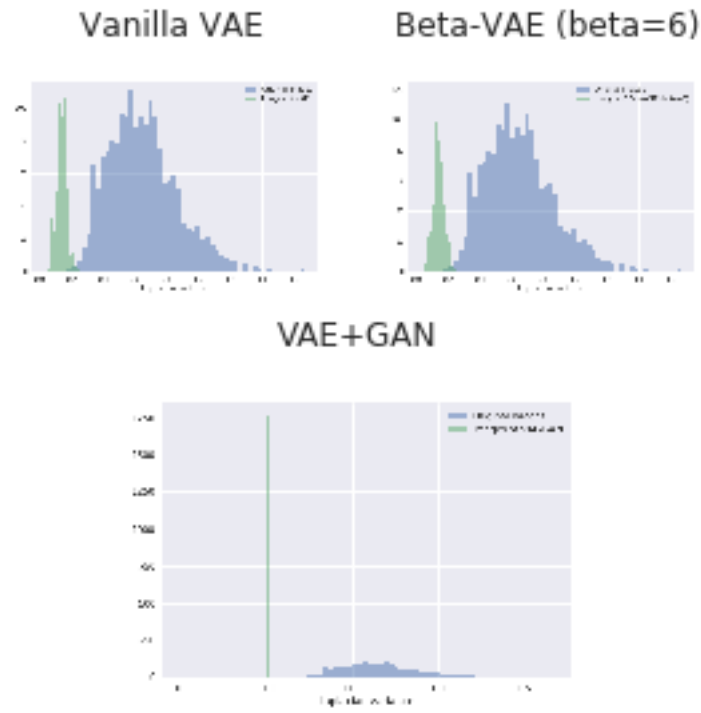
```

In [0]: filename1 = "vae-final/vae_laplacian_variance_reconstructed.png"
        filename2 = "vae-final/bvae_laplacian_variance_reconstructed.png"

```

```
filename3 = "vae-final/vae_gan_laplacian_variance_reconstructed.png"
```

```
display_three_plots(filename1, filename2,filename3,"laplacian_blur.png")
```



0.8 Classifier

```
In [0]: ### Metrics - Base Class For all Metrics
```

```
class Metric:
    def __init__(self):
        pass
    def __call__(self, outputs, target, loss):
        raise NotImplementedError

    def reset(self):
        raise NotImplementedError

    def value(self):
        raise NotImplementedError

    def name(self):
        raise NotImplementedError
```

```

## Accuracy Metric
class AccumulatedAccuracyMetric(Metric):
    def __init__(self):
        self.correct = 0
        self.total = 0

    def __call__(self, outputs, target):
        # Track the accuracy
        _, argmax = torch.max(outputs, 1)
        accuracy = (target == argmax.squeeze()).float().sum()
        self.correct += accuracy
        self.total += target.size(0)
        return self.value()

    def reset(self):
        self.correct = 0
        self.total = 0

    def value(self):
        return 100 * float(self.correct) / self.total

    def name(self):
        return 'Accuracy'

## Loss
class RunningAverage():
    """A simple class that maintains the
    running average of a quantity
    Example:
    ...
    loss_avg = RunningAverage()
    loss_avg.update(2)
    loss_avg.update(4)
    loss_avg() = 3
    ...
    """

    def __init__( self ):
        self.steps = 0
        self.total = 0

    def update( self, val ):
        self.total += val
        self.steps += 1

    def __call__( self ):
        return self.total / float ( self.steps )

```

```

In [0]: class Classifier(nn.Module):
        def __init__(self):
            super(Classifier,self).__init__()

            ## Define NN
            self.fc1 = nn.Linear(10, 10)

        def forward(self,x):
            ## flat input features
            x = x.view(-1, self.num_flat_features(x))
            x = self.fc1(x)

            return F.log_softmax(x, dim=1)

        def num_flat_features(self,x):
            size = x.size()[1:] # all dimensions except the batch dimension
            num_features = 1
            for s in size:
                num_features *=s
            return num_features

In [0]: ## Training
        from tqdm import trange

        criterion = nn.CrossEntropyLoss()
        classifier = Classifier()
        if is_cuda:
            classifier = classifier.to(device)

        # Loss and optimizer
        learning_rate = 0.001
        momentum = 0.9
        criterion = nn.CrossEntropyLoss()
        optimizer = torch.optim.Adam(classifier.parameters(), lr=learning_rate)
        scheduler = lr_scheduler.StepLR(optimizer,step_size = 7,gamma = 0.1)
        print(classifier)

Classifier(
  (fc1): Linear(in_features=10, out_features=10, bias=True)
)

In [0]: def train_classifier_epoch(model, epoch):
        model.train()
        metric = AccumulatedAccuracyMetric()

        for idx, (data, labels) in enumerate(train_loader):

```

```

data= data.to(device)
labels = labels.to(device)

if isinstance(model,ConvVAE):
    recon_batch, mu, logvar = model(data)
else:
    mu, logvar,recon_batch = model(data)

## classifier, pass latent vector
outputs = classifier(mu)
classifier_loss = criterion(outputs, labels)

optimizer.zero_grad()
classifier_loss.backward()
optimizer.step()

classifier_loss /= data.size(0)

metric(outputs, labels)

return 0.0, metric

## Test Epoch
"""
Test, classifier on learnt features
"""
def test_classifier_epoch(model,epoch):
    classifier.eval()
    metric = AccumulatedAccuracyMetric()

    for idx, (data, labels) in enumerate(test_loader):
        data= data.to(device)
        labels = labels.to(device)

        if isinstance(model,ConvVAE):
            recon_batch, mu, logvar = model(data)
        else:
            mu, logvar,recon_batch = model(data)

        ## classifier, pass latent vector
        outputs = classifier(mu)
        classifier_loss = criterion(outputs, labels)

        classifier_loss /= data.size(0)

```

```

        metric(outputs, labels)

    return 0.0, metric

In [0]: ## VAE - vanilla
vae_parameters = list(vae_model.named_parameters())
for name, param in vae_parameters:
    param.requires_grad = False

beta_vae_parameters = list(beta_vae.named_parameters())
for name, param in beta_vae_parameters:
    param.requires_grad = False

vae_gan_parameters = list(G.named_parameters())
for name, param in vae_gan_parameters:
    param.requires_grad = False

In [0]: def train_and_test(model):

    train_accuracy = []

    test_accuracy = []
    n_epochs = 50
    for epoch in range(1, n_epochs):

        # Train stage
        train_loss, metric = train_classifier_epoch(model, epoch)

        train_accuracy.append(metric.value())

        message = 'Epoch: {}/{}. Train set: Average loss: {:.4f}'
            .format(epoch + 1, n_epochs, 0.0)
        message += '\t Average Accuracy: \t{}: {}'.format(
            metric.name(), metric.value())
        print(message)

        val_loss, metrics = test_classifier_epoch(model, epoch)

        test_accuracy.append(metrics.value())

        message += '\nEpoch: {}/{}. Test set: Average loss: {:.4f}'
            .format(epoch + 1, n_epochs, 0.0)

        message += '\t Average Accuracy: \t{}: {}'.format(
            metrics.name(), metrics.value())

```

```
print(message)
```

```
return train_accuracy, test_accuracy
```

```
In [0]: vae_train_ac, vae_test_ac = train_and_test(vae_model)
```

Epoch: 2/50. Train set: Average loss: 0.0058	Average Accuracy:	Accuracy: 88.36
Epoch: 2/50. Train set: Average loss: 0.0058	Average Accuracy:	Accuracy: 88.36
Epoch: 2/50. Test set: Average loss: 0.0049	Average Accuracy:	Accuracy: 92.15
Epoch: 3/50. Train set: Average loss: 0.0038	Average Accuracy:	Accuracy: 92.3
Epoch: 3/50. Train set: Average loss: 0.0038	Average Accuracy:	Accuracy: 92.3
Epoch: 3/50. Test set: Average loss: 0.0034	Average Accuracy:	Accuracy: 93.66
Epoch: 4/50. Train set: Average loss: 0.0029	Average Accuracy:	Accuracy: 93.56
Epoch: 4/50. Train set: Average loss: 0.0029	Average Accuracy:	Accuracy: 93.56
Epoch: 4/50. Test set: Average loss: 0.0028	Average Accuracy:	Accuracy: 94.37
Epoch: 5/50. Train set: Average loss: 0.0024	Average Accuracy:	Accuracy: 94.21
Epoch: 5/50. Train set: Average loss: 0.0024	Average Accuracy:	Accuracy: 94.21
Epoch: 5/50. Test set: Average loss: 0.0027	Average Accuracy:	Accuracy: 94.73
Epoch: 6/50. Train set: Average loss: 0.0021	Average Accuracy:	Accuracy: 94.64
Epoch: 6/50. Train set: Average loss: 0.0021	Average Accuracy:	Accuracy: 94.64
Epoch: 6/50. Test set: Average loss: 0.0022	Average Accuracy:	Accuracy: 95.02
Epoch: 7/50. Train set: Average loss: 0.0019	Average Accuracy:	Accuracy: 94.97
Epoch: 7/50. Train set: Average loss: 0.0019	Average Accuracy:	Accuracy: 94.97
Epoch: 7/50. Test set: Average loss: 0.0019	Average Accuracy:	Accuracy: 95.2
Epoch: 8/50. Train set: Average loss: 0.0018	Average Accuracy:	Accuracy: 95.11
Epoch: 8/50. Train set: Average loss: 0.0018	Average Accuracy:	Accuracy: 95.11
Epoch: 8/50. Test set: Average loss: 0.0020	Average Accuracy:	Accuracy: 95.34
Epoch: 9/50. Train set: Average loss: 0.0016	Average Accuracy:	Accuracy: 95.25
Epoch: 9/50. Train set: Average loss: 0.0016	Average Accuracy:	Accuracy: 95.25
Epoch: 9/50. Test set: Average loss: 0.0018	Average Accuracy:	Accuracy: 95.49
Epoch: 10/50. Train set: Average loss: 0.0016	Average Accuracy:	Accuracy: 95.38
Epoch: 10/50. Train set: Average loss: 0.0016	Average Accuracy:	Accuracy: 95.38
Epoch: 10/50. Test set: Average loss: 0.0015	Average Accuracy:	Accuracy: 95.64
Epoch: 11/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 95.4
Epoch: 11/50. Train set: Average loss: 0.0015	Average Accuracy:	Accuracy: 95.4
Epoch: 11/50. Test set: Average loss: 0.0016	Average Accuracy:	Accuracy: 95.7
Epoch: 12/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.5
Epoch: 12/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.5
Epoch: 12/50. Test set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.76
Epoch: 13/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.6
Epoch: 13/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.6
Epoch: 13/50. Test set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.77
Epoch: 14/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.7
Epoch: 14/50. Train set: Average loss: 0.0014	Average Accuracy:	Accuracy: 95.7
Epoch: 14/50. Test set: Average loss: 0.0013	Average Accuracy:	Accuracy: 95.84
Epoch: 15/50. Train set: Average loss: 0.0013	Average Accuracy:	Accuracy: 95.7
Epoch: 15/50. Train set: Average loss: 0.0013	Average Accuracy:	Accuracy: 95.7
Epoch: 15/50. Test set: Average loss: 0.0013	Average Accuracy:	Accuracy: 95.92

[illegible]


[illegible]

Epoch: 48/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.13
Epoch: 48/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.13
Epoch: 48/50. Test set: Average loss: 0.0013	Average Accuracy:	Accuracy: 96.13
Epoch: 49/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.13
Epoch: 49/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.13
Epoch: 49/50. Test set: Average loss: 0.0011	Average Accuracy:	Accuracy: 96.14
Epoch: 50/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.14
Epoch: 50/50. Train set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.14
Epoch: 50/50. Test set: Average loss: 0.0012	Average Accuracy:	Accuracy: 96.16

```
In [0]: import pickle
def save_lists(l, pickle_filename):
    with open(pickle_filename, 'wb') as f:
        pickle.dump(l, f)
def load_lists(l, picklet_filename):
    with open(pickle_filename, 'rb') as f:
        l = pickle.load(f)

In [0]: ## save lists
save_lists(vae_train_ac, "vae-final/vae-train-classifier-acc.pkl")
save_lists(vae_test_ac, "vae-final/vae-test-classifier-acc.pkl")

In [0]: n_epochs = 50
plt.plot(range(1, n_epochs), vae_train_ac)
plt.plot(range(1, n_epochs), vae_test_ac)
plt.title("Vanilla VAE - max train acc: {}, max test acc: {}%"
          .format(round(max(vae_train_ac), 2), max(vae_test_ac)))
plt.savefig("vae-final/vae-train-test-acc-classifier.png")
```



```
In [0]: bvae_train_ac, bvae_test_ac = train_and_test(beta_vae)
```

Epoch: 2/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 91.91
Epoch: 2/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 91.91
Epoch: 2/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.68
Epoch: 3/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.73
Epoch: 3/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 92.73
Epoch: 3/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.26
Epoch: 4/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.18
Epoch: 4/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.18
Epoch: 4/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.6
Epoch: 5/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.47
Epoch: 5/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.47
Epoch: 5/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.69
Epoch: 6/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.69
Epoch: 6/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.69
Epoch: 6/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.87
Epoch: 7/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.84
Epoch: 7/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.84
Epoch: 7/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.95
Epoch: 8/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.97
Epoch: 8/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 93.97


[illegible]

[illegible]

Epoch: 40/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.74
Epoch: 41/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 41/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 41/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72
Epoch: 42/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 42/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 42/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.71
Epoch: 43/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.70
Epoch: 43/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.70
Epoch: 43/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.7
Epoch: 44/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 44/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 44/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72
Epoch: 45/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 45/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 45/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.76
Epoch: 46/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 46/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 46/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72
Epoch: 47/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 47/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 47/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.71
Epoch: 48/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 48/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 48/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72
Epoch: 49/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.74
Epoch: 49/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.74
Epoch: 49/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72
Epoch: 50/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 50/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.73
Epoch: 50/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 94.72

```
In [0]: ## save lists
        save_lists(bvae_train_ac,"vae-final/bvae-train-classifier-acc.pkl")
        save_lists(bvae_test_ac,"vae-final/bvae-test-classifier-acc.pkl")

        n_epochs = 50
        plt.plot(range(1, n_epochs),bvae_train_ac)
        plt.plot(range(1,n_epochs),bvae_test_ac)
        plt.title("Beta VAE - max train acc: {}%, max test acc: {}%"
                  .format(round(max(bvae_train_ac),2),max(bvae_test_ac)))
        plt.savefig("vae-final/bvae-train-test-acc-classifier.png")
```



```
In [0]: vae_gan_train_ac, vae_gan_test_ac = train_and_test(G)
```

Epoch: 2/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 26.76
Epoch: 2/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 26.76
Epoch: 2/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 33.57
Epoch: 3/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 39.71
Epoch: 3/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 39.71
Epoch: 3/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 47.65
Epoch: 4/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 53.27
Epoch: 4/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 53.27
Epoch: 4/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 61.43
Epoch: 5/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 65.35
Epoch: 5/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 65.35
Epoch: 5/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 72.0
Epoch: 6/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 74.37
Epoch: 6/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 74.37
Epoch: 6/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 78.51
Epoch: 7/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 79.64
Epoch: 7/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 79.64
Epoch: 7/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 82.52
Epoch: 8/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 82.80
Epoch: 8/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 82.80


[illegible]

[illegible]

Epoch: 40/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.78
Epoch: 41/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 41/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 41/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.8
Epoch: 42/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 42/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 42/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.73
Epoch: 43/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 43/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 43/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.74
Epoch: 44/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 44/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 44/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.69
Epoch: 45/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 45/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 45/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.81
Epoch: 46/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 46/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 46/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.84
Epoch: 47/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 47/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 47/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.73
Epoch: 48/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 48/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 48/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.79
Epoch: 49/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 49/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.23
Epoch: 49/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.8
Epoch: 50/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.3
Epoch: 50/50. Train set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.3
Epoch: 50/50. Test set: Average loss: 0.0000	Average Accuracy:	Accuracy: 89.81

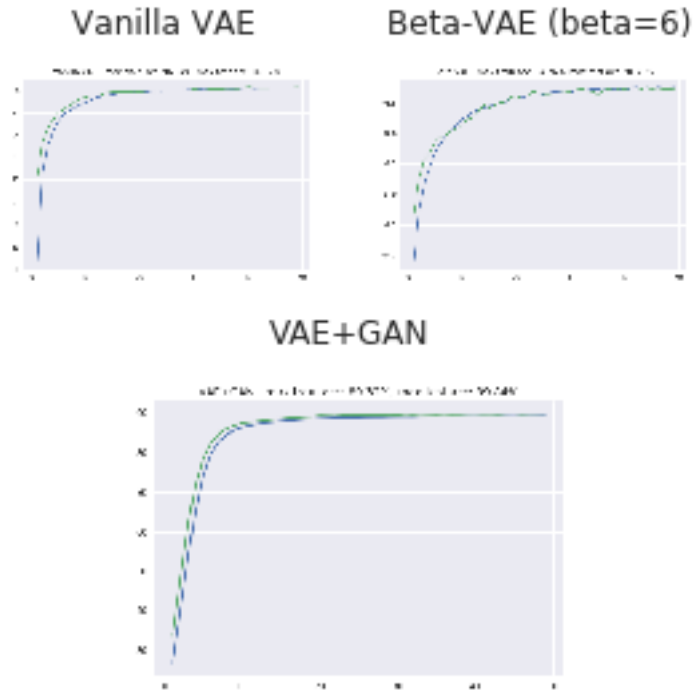
```
In [0]: ## save lists
        save_lists(vae_gan_train_ac,"vae-final/vaeg-train-classifier-acc.pkl")
        save_lists(vae_gan_test_ac,"vae-final/vaeg-test-classifier-acc.pkl")

        n_epochs = 50
        plt.plot(range(1, n_epochs),vae_gan_train_ac)
        plt.plot(range(1,n_epochs),vae_gan_test_ac)
        plt.title("VAE+GAN - max train acc: {}%, max test acc: {}%"
                  .format(round(max(vae_gan_train_ac),2),max(vae_gan_test_ac)))
        plt.savefig("vae-final/vae_gan-train-test-acc-classifier.png")
```



```
In [0]: import matplotlib.pyplot as plt
        from IPython.display import Image
        import matplotlib.image as mpimg
        import matplotlib.gridspec as gridspec
        filename1 = "vae-final/vae-train-test-acc-classifier.png"
        filename2 = "vae-final/bvae-train-test-acc-classifier.png"
        filename3 = "vae-final/vae_gan-train-test-acc-classifier.png"

        display_three_plots(filename1, filename2,filename3,"classifier.png")
```

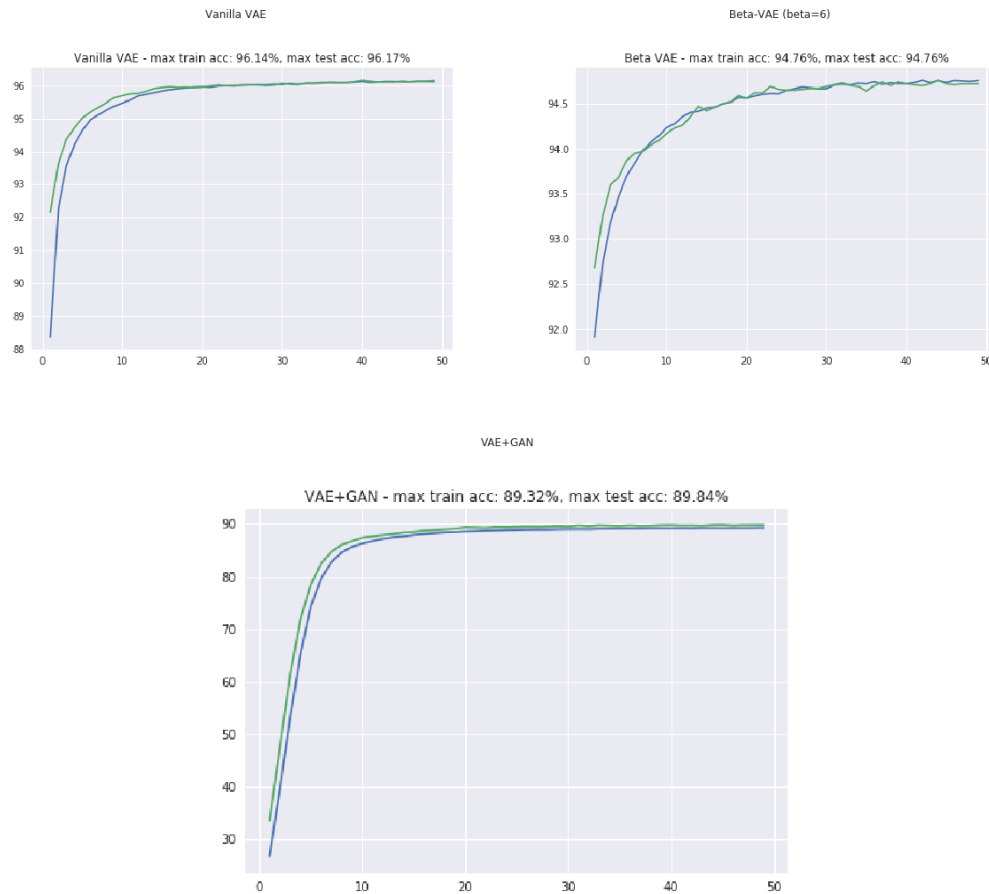


```
In [0]: fig = plt.figure(figsize = (22,20))
gs = gridspec.GridSpec(2,2)
gs.update(wspace=0.0, hspace=0.0)
plt.grid(b=None)
## Train
ax1 = plt.subplot(gs[0, 0])
vae_tsne = filename1
img=mpimg.imread(vae_tsne)
imgplot = plt.imshow(img)
plt.title("Vanilla VAE")
plt.axis("off")
ax1.set_aspect('equal')
ax2 = plt.subplot(gs[0, 1])
betavae_tsne = filename2
img=mpimg.imread(betavae_tsne)
imgplot = plt.imshow(img)
plt.title("Beta-VAE (beta=6)")
plt.axis("off")
ax2.set_aspect('equal')
ax3 = plt.subplot(gs[1,:])
vae_gan_tsne =filename3
img=mpimg.imread(vae_gan_tsne)
```

```

imgplot = plt.imshow(img)
ax3.set_aspect('equal')
plt.title("VAE+GAN")
plt.axis("off")
plt.savefig("vae-final/classifier.png")

```



```

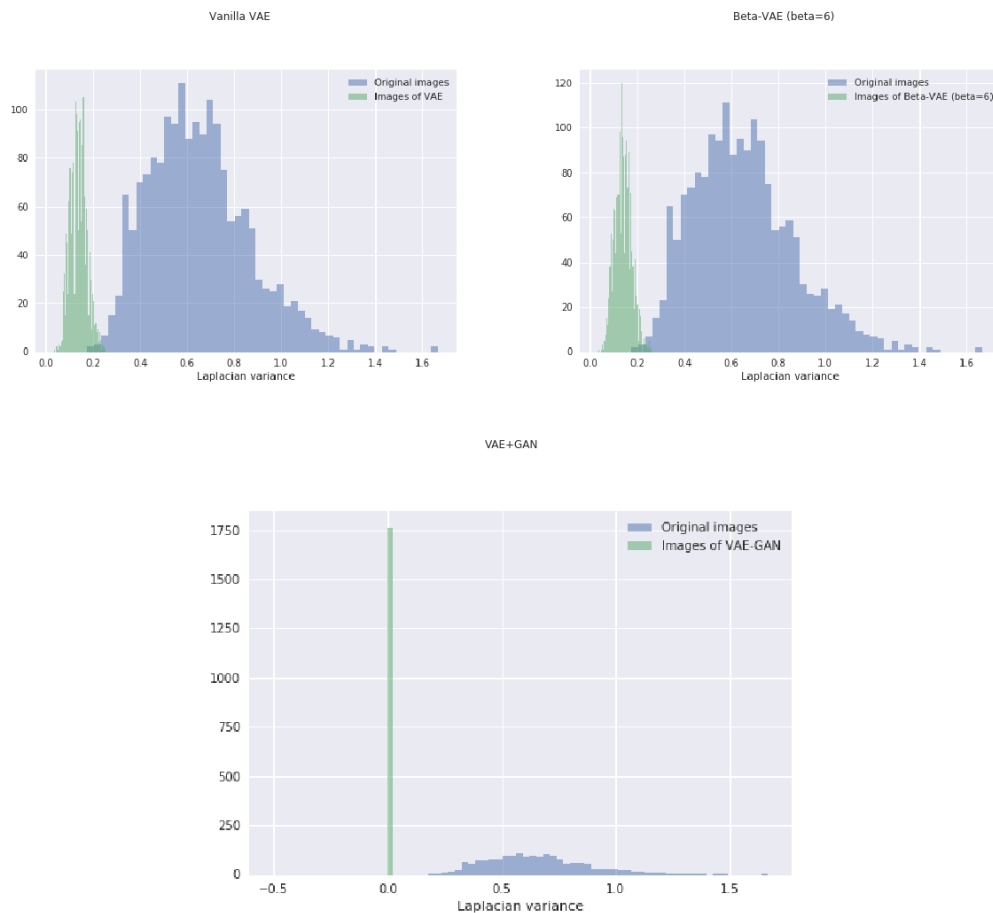
In [0]: fig = plt.figure(figsize = (22,20))
gs = gridspec.GridSpec(2,2)
gs.update(wspace=0.0, hspace=0.0)
plt.grid(b=None)
## Train
ax1 = plt.subplot(gs[0, 0])
vae_tsne = filename1
img=mpimg.imread(vae_tsne)
imgplot = plt.imshow(img)
plt.title("Vanilla VAE")
plt.axis("off")
ax1.set_aspect('equal')

```

```

ax2 = plt.subplot(gs[0, 1])
betavae_tsne = filename2
img=mpimg.imread(betavae_tsne)
imgplot = plt.imshow(img)
plt.title("Beta-VAE (beta=6)")
plt.axis("off")
ax2.set_aspect('equal')
ax3 = plt.subplot(gs[1,:])
vae_gan_tsne =filename3
img=mpimg.imread(vae_gan_tsne)
imgplot = plt.imshow(img)
ax3.set_aspect('equal')
plt.title("VAE+GAN")
plt.axis("off")
plt.savefig("vae-final/laplacian_blur.png")

```



In [0]: